



☰ Fundamentals > Agent Loop

Fundamentals

Agent Loop

An agentic loop is the full “real” run of an agent: intake → context assembly → model inference → tool execution → streaming replies → persistence. It’s the authoritative path that turns a message into actions and a final reply, while keeping session state consistent.

In OpenClaw, a loop is a single, serialized run per session that emits lifecycle and stream events as the model thinks, calls tools, and streams output. This doc explains how that authentic loop is wired end-to-end.

Entry points

Gateway RPC: `agent` and `agent.wait`.

CLI: `agent command`.

How it works (high-level)

1. `agent` RPC validates params, resolves session (`sessionKey/sessionId`), persists session metadata, returns `{ runId, acceptedAt }` immediately.
2. `agentCommand` runs the agent:
 - resolves model + thinking/verbose defaults
 - loads skills snapshot



calls `runEmbeddedPiAgent` (pi-agent-core runtime)
emits **lifecycle end/error** if the embedded loop does not emit one
>

3. `runEmbeddedPiAgent` :

serializes runs via per-session + global queues
resolves model + auth profile and builds the pi session
subscribes to pi events and streams assistant/tool deltas
enforces timeout → aborts run if exceeded
returns payloads + usage metadata

4. `subscribeEmbeddedPiSession` bridges pi-agent-core events to OpenClaw agent stream:

tool events ⇒ stream: "tool"
assistant deltas ⇒ stream: "assistant"
lifecycle events ⇒ stream: "lifecycle" (phase: "start" | "end" | "error")

5. `agent.wait` uses `waitForAgentJob` :

waits for **lifecycle end/error** for `runId`
returns { status: ok|error|timeout, startedAt, endedAt, error? }

Queueing + concurrency

Runs are serialized per session key (session lane) and optionally through a global lane.

This prevents tool/session races and keeps session history consistent.

Messaging channels can choose queue modes (collect/steer/followup) that feed this lane system. See [Command Queue](#).

Session + workspace preparation



Workspace is resolved and created; sandboxed runs may redirect to a sandbox workspace root.

Skills are loaded (or reused from a snapshot) and injected into env and prompt.

Bootstrap/context files are resolved and injected into the system prompt report.

A session write lock is acquired; SessionManager is opened and prepared before streaming.

Prompt assembly + system prompt

System prompt is built from OpenClaw's base prompt, skills prompt, bootstrap context, and per-run overrides.

Model-specific limits and compaction reserve tokens are enforced.

See [System prompt](#) for what the model sees.

Hook points (where you can intercept)

OpenClaw has two hook systems:

Internal hooks (Gateway hooks): event-driven scripts for commands and lifecycle events.

Plugin hooks: extension points inside the agent/tool lifecycle and gateway pipeline.

Internal hooks (Gateway hooks)

`agent:bootstrap` : runs while building bootstrap files before the system prompt is finalized. Use this to add/remove bootstrap context files.

 **Command hooks:** /new , /reset , /stop , and other command events (see Hooks doc).

See [Hooks](#) for setup and examples.

Plugin hooks (agent + gateway lifecycle)

These run inside the agent loop or gateway pipeline:

```
before_model_resolve : runs pre-session (no messages) to  
deterministically override provider/model before model resolution.  
  
before_prompt_build : runs after session load (with messages) to  
inject prependContext / systemPrompt before prompt submission.  
  
before_agent_start : legacy compatibility hook that may run in either  
phase; prefer the explicit hooks above.  
  
agent_end : inspect the final message list and run metadata after  
completion.  
  
before_compaction / after_compaction : observe or annotate compaction  
cycles.  
  
before_tool_call / after_tool_call : intercept tool params/results.  
  
tool_result_persist : synchronously transform tool results before  
they are written to the session transcript.  
  
message_received / message_sending / message_sent : inbound + outbound  
message hooks.  
  
session_start / session_end : session lifecycle boundaries.  
  
gateway_start / gateway_stop : gateway lifecycle events.
```

See [Plugins](#) for the hook API and registration details.

Streaming + partial replies



Assistant deltas are streamed from pi-agent-core and emitted as assistant events.

Block streaming can emit partial replies either on `text_end` or `message_end`.

Reasoning streaming can be emitted as a separate stream or as block replies.

See [**Streaming**](#) for chunking and block reply behavior.

Tool execution + messaging tools

Tool start/update/end events are emitted on the `tool` stream.

Tool results are sanitized for size and image payloads before logging/emitting.

Messaging tool sends are tracked to suppress duplicate assistant confirmations.

Reply shaping + suppression

Final payloads are assembled from:

assistant text (and optional reasoning)

inline tool summaries (when verbose + allowed)

assistant error text when the model errors

`NO_REPLY` is treated as a silent token and filtered from outgoing payloads.

Messaging tool duplicates are removed from the final payload list.

If no renderable payloads remain and a tool errored, a fallback tool error reply is emitted (unless a messaging tool already sent a user-visible reply).

Compaction + retries



Auto-compaction emits `compaction` stream events and can trigger a retry. >

On retry, in-memory buffers and tool summaries are reset to avoid duplicate output.

See [Compaction](#) for the compaction pipeline.

Event streams (today)

`lifecycle` : emitted by `subscribeEmbeddedPiSession` (and as a fallback by `agentCommand`)

`assistant` : streamed deltas from pi-agent-core

`tool` : streamed tool events from pi-agent-core

Chat channel handling

Assistant deltas are buffered into chat `delta` messages.

A chat `final` is emitted on `lifecycle end/error`.

Timeouts

`agent.wait default: 30s` (just the wait). `timeoutMs` param overrides.

Agent runtime: `agents.defaults.timeoutSeconds default 600s`; enforced in `runEmbeddedPiAgent` abort timer.

Where things can end early

Agent timeout (abort)

AbortSignal (cancel)

Gateway disconnect or RPC timeout

```
agent.wait timeout (wait-only, does not stop agent)
```



[« Agent Runtime »](#)

[System Prompt »](#)

Powered by [mintlify](#)