



☰ Sessions and memory > Compaction

Sessions and memory

Compaction

Every model has a **context window** (max tokens it can see). Long-running chats accumulate messages and tool results; once the window is tight, OpenClaw **compacts** older history to stay within limits.

What compaction is

Compaction **summarizes older conversation** into a compact summary entry and keeps recent messages intact. The summary is stored in the session history, so future requests use:

The compaction summary

Recent messages after the compaction point

Compaction **persists** in the session's JSONL history.

Configuration

Use the `agents.defaults.compaction` setting in your `openclaw.json` to configure compaction behavior (mode, target tokens, etc.).

Auto-compaction (default on)

When a session nears or exceeds the model's context window, OpenClaw triggers auto-compaction and may retry the original request using the

compacted context.



You'll see:

```
>
  ✎ Auto-compaction complete in verbose mode
  /status showing ✎ Compactions: <count>
```

Before compaction, OpenClaw can run a **silent memory flush** turn to store durable notes to disk. See [Memory](#) for details and config.

Manual compaction

Use `/compact` (optionally with instructions) to force a compaction pass:

```
/compact Focus on decisions and open questions
```

Context window source

Context window is model-specific. OpenClaw uses the model definition from the configured provider catalog to determine limits.

Compaction vs pruning

Compaction: summarises and **persists** in JSONL.

Session pruning: trims old **tool results** only, **in-memory**, per request.

See [Pruning](#) for pruning details.

Tips

Use `/compact` when sessions feel stale or context is bloated.



Large tool outputs are already truncated; pruning can further reduce tool-result buildup.

If you need a fresh slate, `/new` or `/reset` starts a new session id.

< Memory

Multi-Agent Routing >

Powered by [mintlify](#)