



☰ Fundamentals > Context

Fundamentals

Context

“Context” is **everything OpenClaw sends to the model for a run**. It is bounded by the model’s **context window** (token limit).

Beginner mental model:

System prompt (OpenClaw-built): rules, tools, skills list, time/runtime, and injected workspace files.

Conversation history: your messages + the assistant’s messages for this session.

Tool calls/results + attachments: command output, file reads, images/audio, etc.

Context is *not the same thing* as “memory”: memory can be stored on disk and reloaded later; context is what’s inside the model’s current window.

Quick start (inspect context)

`/status` → quick “how full is my window?” view + session settings.

`/context list` → what’s injected + rough sizes (per file + totals).

`/context detail` → deeper breakdown: per-file, per-tool schema sizes, per-skill entry sizes, and system prompt size.

`/usage tokens` → append per-reply usage footer to normal replies.

 **/compact** → summarize older history into a compact entry to free window space.

See also: [Slash commands](#), [Token use & costs](#), [Compaction](#).

Example output

Values vary by model, provider, tool policy, and what's in your workspace.

/context list

 Context breakdown

Workspace: <workspaceDir>

Bootstrap max/file: 20,000 chars

Sandbox: mode=non-main sandboxed=false

System prompt (run): 38,412 chars (~9,603 tok) (Project Context 23,901 chars (~5,100 tok))

Injected workspace files:

- AGENTS.md: OK | raw 1,742 chars (~436 tok) | injected 1,742 chars (~436 tok)
- SOUL.md: OK | raw 912 chars (~228 tok) | injected 912 chars (~228 tok)
- TOOLS.md: TRUNCATED | raw 54,210 chars (~13,553 tok) | injected 20,962 chars (~5,100 tok)
- IDENTITY.md: OK | raw 211 chars (~53 tok) | injected 211 chars (~53 tok)
- USER.md: OK | raw 388 chars (~97 tok) | injected 388 chars (~97 tok)
- HEARTBEAT.md: MISSING | raw 0 | injected 0
- BOOTSTRAP.md: OK | raw 0 chars (~0 tok) | injected 0 chars (~0 tok)

Skills list (system prompt text): 2,184 chars (~546 tok) (12 skills)

Tools: read, edit, write, exec, process, browser, message, sessions_send, ...

Tool list (system prompt text): 1,032 chars (~258 tok)

Tool schemas (JSON): 31,988 chars (~7,997 tok) (counts toward context; not shown above)

Tools: (same as above)

Session tokens (cached): 14,250 total / ctx=32,000

/context detail



🧠 Context breakdown (detailed)

>

...

Top skills (prompt entry size):

- frontend-design: 412 chars (~103 tok)
- oracle: 401 chars (~101 tok)
- ... (+10 more skills)

Top tools (schema size):

- browser: 9,812 chars (~2,453 tok)
- exec: 6,240 chars (~1,560 tok)
- ... (+N more tools)

What counts toward the context window

Everything the model receives counts, including:

System prompt (all sections).

Conversation history.

Tool calls + tool results.

Attachments/transcripts (images/audio/files).

Compaction summaries and pruning artifacts.

Provider “wrappers” or hidden headers (not visible, still counted).

How OpenClaw builds the system prompt

The system prompt is **OpenClaw-owned** and rebuilt each run. It includes:

Tool list + short descriptions.

Skills list (metadata only; see below).

Workspace location.

Time (UTC + converted user time if configured).



Runtime metadata (host/OS/model/thinking).

Injected workspace bootstrap files under **Project Context**.

Full breakdown: [System Prompt](#).

Injected workspace files (Project Context)

By default, OpenClaw injects a fixed set of workspace files (if present):

AGENTS.md

SOUL.md

TOOLS.md

IDENTITY.md

USER.md

HEARTBEAT.md

BOOTSTRAP.md (first-run only)

Large files are truncated per-file using `agents.defaults.bootstrapMaxChars` (default 20000 chars). OpenClaw also enforces a total bootstrap injection cap across files with `agents.defaults.bootstrapTotalMaxChars` (default 150000 chars). `/context` shows **raw vs injected** sizes and whether truncation happened.

Skills: what's injected vs loaded on-demand

The system prompt includes a compact **skills list** (name + description + location). This list has real overhead.

Skill instructions are *not* included by default. The model is expected to read the skill's SKILL.md only when needed.

Tools: there are two costs



Tools affect context in two ways:

- >
- 1. **Tool list text** in the system prompt (what you see as “Tooling”).
- 2. **Tool schemas** (JSON). These are sent to the model so it can call tools. They count toward context even though you don’t see them as plain text.

`/context detail` breaks down the biggest tool schemas so you can see what dominates.

Commands, directives, and “inline shortcuts”

Slash commands are handled by the Gateway. There are a few different behaviors:

Standalone commands: a message that is only `/...` runs as a command.

Directives: `/think` , `/verbose` , `/reasoning` , `/elevated` , `/model` , `/queue` are stripped before the model sees the message.

Directive-only messages persist session settings.

Inline directives in a normal message act as per-message hints.

Inline shortcuts (allowlisted senders only): certain `/...` tokens inside a normal message can run immediately (example: “hey `/status`”), and are stripped before the model sees the remaining text.

Details: [Slash commands](#).

Sessions, compaction, and pruning (what persists)

What persists across messages depends on the mechanism:



Normal history persists in the session transcript until compacted/pruned by policy.

Compaction persists a summary into the transcript and keeps recent messages intact.

Pruning removes old tool results from the *in-memory* prompt for a run, but does not rewrite the transcript.

Docs: [Session](#), [Compaction](#), [Session pruning](#).

What /context actually reports

/context prefers the latest **run-built** system prompt report when available:

System prompt (run) = captured from the last embedded (tool-capable) run and persisted in the session store.

System prompt (estimate) = computed on the fly when no run report exists (or when running via a CLI backend that doesn't generate the report).

Either way, it reports sizes and top contributors; it does **not** dump the full system prompt or tool schemas.

[**< System Prompt**](#)

[**Agent Workspace >**](#)

Powered by [mintlify](#)