



### ☰ Security and sandboxing > **Sandbox vs Tool Policy vs Elevated**

#### Security and sandboxing

## Sandbox vs Tool Policy vs Elevated

OpenClaw has three related (but different) controls:

1. **Sandbox** (`agents.defaults.sandbox.*` / `agents.list[].sandbox.*`) decides **where tools run** (Docker vs host).
2. **Tool policy** (`tools.*`, `tools.sandbox.tools.*`, `agents.list[].tools.*`) decides **which tools are available/allowed**.
3. **Elevated** (`tools.elevated.*`, `agents.list[].tools.elevated.*`) is an **exec-only escape hatch** to run on the host when you're sandboxed.

## Quick debug

Use the inspector to see what OpenClaw is *actually* doing:

```
openclaw sandbox explain  
openclaw sandbox explain --session agent:main:main  
openclaw sandbox explain --agent work  
openclaw sandbox explain --json
```

It prints:

```
effective sandbox mode/scope/workspace access  
whether the session is currently sandboxed (main vs non-main)  
effective sandbox tool allow/deny (and whether it came from  
agent/global/default)
```

elevated gates and fix-it key paths



## Sandbox: where tools run

Sandboxing is controlled by `agents.defaults.sandbox.mode` :

"off" : everything runs on the host.

"non-main" : only non-main sessions are sandboxed (common "surprise" for groups/channels).

"all" : everything is sandboxed.

See [Sandboxing](#) for the full matrix (scope, workspace mounts, images).

## Bind mounts (security quick check)

`docker.binds` pierces the sandbox filesystem: whatever you mount is visible inside the container with the mode you set ( `:ro` or `:rw` ).

Default is read-write if you omit the mode; prefer `:ro` for source/secrets.

`scope: "shared"` ignores per-agent binds (only global binds apply).

Binding `/var/run/docker.sock` effectively hands host control to the sandbox; only do this intentionally.

Workspace access (`workspaceAccess: "ro" / "rw"`) is independent of bind modes.

## Tool policy: which tools exist/are callable

Two layers matter:

**Tool profile:** `tools.profile` and `agents.list[].tools.profile` (base allowlist)



**Provider tool profile:** tools.byProvider[provider].profile and  
agents.list[].tools.byProvider[provider].profile

**Global/per-agent tool policy:** tools.allow / tools.deny and  
agents.list[].tools.allow / agents.list[].tools.deny

**Provider tool policy:** tools.byProvider[provider].allow/deny and  
agents.list[].tools.byProvider[provider].allow/deny

**Sandbox tool policy** (only applies when sandboxed):  
tools.sandbox.tools.allow / tools.sandbox.tools.deny and  
agents.list[].tools.sandbox.tools.\*

Rules of thumb:

deny always wins.

If allow is non-empty, everything else is treated as blocked.

Tool policy is the hard stop: /exec cannot override a denied exec tool.

/exec only changes session defaults for authorized senders; it does not grant tool access. Provider tool keys accept either provider (e.g. google-antigravity) or provider/model (e.g. openai/gpt-5.2).

## Tool groups (shorthands)

Tool policies (global, agent, sandbox) support group:\* entries that expand to multiple tools:



```
tools: {
    sandbox: {
        tools: {
            allow: ["group:runtime", "group:fs", "group:sessions", "group:memory"],
        },
    },
},
}
```

Available groups:

```
group:runtime : exec , bash , process
group:fs : read , write , edit , apply_patch
group:sessions : sessions_list , sessions_history , sessions_send ,
sessions_spawn , session_status
group:memory : memory_search , memory_get
group:ui : browser , canvas
group:automation : cron , gateway
group:messaging : message
group:nodes : nodes
group:openclaw : all built-in OpenClaw tools (excludes provider
plugins)
```

## Elevated: exec-only “run on host”

Elevated does **not** grant extra tools; it only affects `exec`.

If you’re sandboxed, `/elevated on` (or `exec` with `elevated: true`) runs on the host (approvals may still apply).

Use `/elevated full` to skip `exec` approvals for the session.

 If you're already running direct, elevated is effectively a no-op (still gated).

Elevated is **not** skill-scoped and does **not** override tool allow/deny.

`/exec` is separate from elevated. It only adjusts per-session exec defaults for authorized senders.

Gates:

Enablement: `tools.elevated.enabled` (and optionally `agents.list[].tools.elevated.enabled`)

Sender allowlists: `tools.elevated.allowFrom.<provider>` (and optionally `agents.list[].tools.elevated.allowFrom.<provider>`)

See [Elevated Mode](#).

## Common “sandbox jail” fixes

### “Tool X blocked by sandbox tool policy”

Fix-it keys (pick one):

Disable sandbox: `agents.defaults.sandbox.mode=off` (or per-agent `agents.list[].sandbox.mode=off`)

Allow the tool inside sandbox:

remove it from `tools.sandbox.tools.deny` (or per-agent `agents.list[].tools.sandbox.tools.deny`)

or add it to `tools.sandbox.tools.allow` (or per-agent allow)

### “I thought this was main, why is it sandboxed?”

In “non-main” mode, group/channel keys are *not* main. Use the main session key (shown by `sandbox explain`) or switch mode to “off”.

[\*\*< Sandboxing\*\*](#)[\*\*Gateway Protocol >\*\*](#)Powered by **mintlify**