



## ☰ Security and sandboxing › Sandboxing

### Security and sandboxing

## Sandboxing

OpenClaw can run **tools inside Docker containers** to reduce blast radius. This is **optional** and controlled by configuration (`agents.defaults.sandbox` or `agents.list[].sandbox`). If sandboxing is off, tools run on the host. The Gateway stays on the host; tool execution runs in an isolated sandbox when enabled.

This is not a perfect security boundary, but it materially limits filesystem and process access when the model does something dumb.

## What gets sandboxed

Tool execution (`exec`, `read`, `write`, `edit`, `apply_patch`, `process`, etc.).

Optional sandboxed browser (`agents.defaults.sandbox.browser`).

By default, the sandbox browser auto-starts (ensures CDP is reachable) when the browser tool needs it. Configure via `agents.defaults.sandbox.browser.autoStart` and `agents.defaults.sandbox.browser.autoStartTimeoutMs`.

`agents.defaults.sandbox.browser.allowHostControl` lets sandboxed sessions target the host browser explicitly.

Optional allowlists gate `target: "custom" : allowedControlUrls`, `allowedControlHosts`, `allowedControlPorts`.

Not sandboxed:

The Gateway process itself.

 Any tool explicitly allowed to run on the host (e.g. tools.elevated ).

**Elevated exec runs on the host and bypasses sandboxing.**

If sandboxing is off, tools.elevated does not change execution (already on host). See [Elevated Mode](#).

## Modes

`agents.defaults.sandbox.mode` controls **when** sandboxing is used:

"off" : no sandboxing.

"non-main" : sandbox only **non-main** sessions (default if you want normal chats on host).

"all" : every session runs in a sandbox. Note: "non-main" is based on `session.mainKey` (default "main"), not agent id. Group/channel sessions use their own keys, so they count as non-main and will be sandboxed.

## Scope

`agents.defaults.sandbox.scope` controls **how many containers** are created:

"session" (default): one container per session.

"agent" : one container per agent.

"shared" : one container shared by all sandboxed sessions.

## Workspace access

`agents.defaults.sandbox.workspaceAccess` controls **what the sandbox can see**:



"none" (default): tools see a sandbox workspace under  
~/.openclaw/sandboxes .

"ro" : mounts the agent workspace read-only at /agent (disables  
write / edit / apply\_patch ).

"rw" : mounts the agent workspace read/write at /workspace .

Inbound media is copied into the active sandbox workspace  
( media/inbound/\* ). Skills note: the read tool is sandbox-rooted. With  
workspaceAccess: "none" , OpenClaw mirrors eligible skills into the  
sandbox workspace ( .../skills ) so they can be read. With "rw" ,  
workspace skills are readable from /workspace/skills .

## Custom bind mounts

agents.defaults.sandbox.docker.binds mounts additional host directories  
into the container. Format: host:container:mode (e.g.,  
"/home/user/source:/source:rw" ).

Global and per-agent binds are merged (not replaced). Under scope:  
"shared" , per-agent binds are ignored.

agents.defaults.sandbox.browser.binds mounts additional host directories  
into the **sandbox browser** container only.

When set (including [] ), it replaces  
agents.defaults.sandbox.docker.binds for the browser container.

When omitted, the browser container falls back to  
agents.defaults.sandbox.docker.binds (backwards compatible).

Example (read-only source + an extra data directory):



```
agents: {
    defaults: {
        sandbox: {           >
            docker: {
                binds: ["/home/user/source:/source:ro", "/var/data/myapp:/data:ro"],
            },
        },
    },
    list: [
        {
            id: "build",
            sandbox: {
                docker: {
                    binds: ["/mnt/cache:/cache:rw"],
                },
            },
        },
    ],
},
}
```

## Security notes:

Binds bypass the sandbox filesystem: they expose host paths with whatever mode you set ( `:ro` or `:rw` ).

OpenClaw blocks dangerous bind sources (for example: `docker.sock` , `/etc` , `/proc` , `/sys` , `/dev` , and parent mounts that would expose them) .

Sensitive mounts (secrets, SSH keys, service credentials) should be `:ro` unless absolutely required.

Combine with `workspaceAccess: "ro"` if you only need read access to the workspace; bind modes stay independent.

See [tool policy](#) and [elevated exec](#) for how binds interact with

## Images + setup

Default image: `openclaw-sandbox:bookworm-slim`

Build it once:

```
scripts/sandbox-setup.sh
```

Note: the default image does **not** include Node. If a skill needs Node (or other runtimes), either bake a custom image or install via `sandbox.docker.setupCommand` (requires network egress + writable root + root user).

Sandboxed browser image:

```
scripts/sandbox-browser-setup.sh
```

By default, sandbox containers run with **no network**. Override with `agents.defaults.sandbox.docker.network`.

Docker installs and the containerized gateway live here:

## setupCommand (one-time container setup)

`setupCommand` runs **once** after the sandbox container is created (not on every run). It executes inside the container via `sh -lc`.

Paths:

Global: `agents.defaults.sandbox.docker.setupCommand`

Per-agent: `agents.list[].sandbox.docker.setupCommand`

Common pitfalls:

 Default docker.network is "none" (no egress), so package installs will fail.

`readOnlyRoot: true` prevents writes; set `readOnlyRoot: false` or bake a custom image.

`user` must be root for package installs (omit `user` or set `user: "0:0"`).

Sandbox exec does **not** inherit host `process.env`. Use `agents.defaults.sandbox.docker.env` (or a custom image) for skill API keys.

## Tool policy + escape hatches

Tool allow/deny policies still apply before sandbox rules. If a tool is denied globally or per-agent, sandboxing doesn't bring it back.

`tools.elevated` is an explicit escape hatch that runs `exec` on the host. `/exec` directives only apply for authorized senders and persist per session; to hard-disable `exec`, use tool policy deny (see [Sandbox vs Tool Policy vs Elevated](#)).

Debugging:

Use `openclaw sandbox explain` to inspect effective sandbox mode, tool policy, and fix-it config keys.

See [Sandbox vs Tool Policy vs Elevated](#) for the “why is this blocked?” mental model. Keep it locked down.

## Multi-agent overrides

Each agent can override sandbox + tools: `agents.list[].sandbox` and `agents.list[].tools` (plus `agents.list[].tools.sandbox.tools` for sandbox tool policy). See [Multi-Agent Sandbox & Tools](#) for precedence.

# Minimal enable example

```
{  
  agents: {  
    defaults: {  
      sandbox: {  
        mode: "non-main",  
        scope: "session",  
        workspaceAccess: "none",  
      },  
    },  
  },  
}
```

## Related docs

[◀ Security](#)[Sandbox vs Tool Policy vs Elevated ▶](#)

Powered by [mintlify](#)