



Security and sandboxing

Security

Quick check: openclaw security audit

See also: [Formal Verification \(Security Models\)](#)

Run this regularly (especially after changing config or exposing network surfaces):

```
openclaw security audit
openclaw security audit --deep
openclaw security audit --fix
openclaw security audit --json
```

It flags common footguns (Gateway auth exposure, browser control exposure, elevated allowlists, filesystem permissions).

OpenClaw is both a product and an experiment: you're wiring frontier-model behavior into real messaging surfaces and real tools. **There is no “perfectly secure” setup.** The goal is to be deliberate about:

- who can talk to your bot
- where the bot is allowed to act
- what the bot can touch

Start with the smallest access that still works, then widen it as you gain confidence.

Hardened baseline in 60 seconds



Use this baseline first, then selectively re-enable tools per trusted agent:

```
{
  gateway: {
    mode: "local",
    bind: "loopback",
    auth: { mode: "token", token: "replace-with-long-random-token" },
  },
  session: {
    dmScope: "per-channel-peer",
  },
  tools: {
    profile: "messaging",
    deny: ["group:automation", "group:runtime", "group:fs", "sessions_spawn", "se
      fs: { workspaceOnly: true },
      exec: { security: "deny", ask: "always" },
      elevated: { enabled: false },
    },
    channels: {
      whatsapp: { dmPolicy: "pairing", groups: { "*": { requireMention: true } } },
    },
  }
}
```

This keeps the Gateway local-only, isolates DMs, and disables control-plane/runtime tools by default.

Shared inbox quick rule

If more than one person can DM your bot:

Set `session.dmScope: "per-channel-peer"` (or `"per-account-channel-peer"` for multi-account channels).

Keep `dmPolicy: "pairing"` or strict allowlists.



Never combine shared DMs with broad tool access.

What the audit checks (high level)

Inbound access (DM policies, group policies, allowlists): can strangers trigger the bot?

Tool blast radius (elevated tools + open rooms): could prompt injection turn into shell/file/network actions?

Network exposure (Gateway bind/auth, Tailscale Serve/Funnel, weak/short auth tokens).

Browser control exposure (remote nodes, relay ports, remote CDP endpoints).

Local disk hygiene (permissions, symlinks, config includes, “synced folder” paths).

Plugins (extensions exist without an explicit allowlist).

Policy drift/misconfig (sandbox docker settings configured but sandbox mode off; ineffective `gateway.nodes.denyCommands` patterns; global `tools.profile="minimal"` overridden by per-agent profiles; extension plugin tools reachable under permissive tool policy).

Model hygiene (warn when configured models look legacy; not a hard block).

If you run `--deep`, OpenClaw also attempts a best-effort live Gateway probe.

Credential storage map

Use this when auditing access or deciding what to back up:

WhatsApp: `~/.openclaw/credentials/whatsapp/<accountId>/creds.json`

Telegram bot token: config/env or `channels.telegram.tokenFile`

Discord bot token: config/env (token file not yet supported)

-  **Slack tokens:** config/env (channels.slack.*)
 -  **Pairing allowlists:** ~/.openclaw/credentials/<channel>-allowFrom.json
 - Model auth profiles:** ,~/.openclaw/agents/<agentId>/agent/auth-profiles.json
 - Legacy OAuth import:** ~/.openclaw/credentials/oauth.json
-

Security Audit Checklist

When the audit prints findings, treat this as a priority order:

1. **Anything “open” + tools enabled:** lock down DMs/groups first (pairing/allowlists), then tighten tool policy/sandboxing.
2. **Public network exposure** (LAN bind, Funnel, missing auth): fix immediately.
3. **Browser control remote exposure:** treat it like operator access (tailnet-only, pair nodes deliberately, avoid public exposure).
4. **Permissions:** make sure state/config/credentials/auth are not group/world-readable.
5. **Plugins/extensions:** only load what you explicitly trust.
6. **Model choice:** prefer modern, instruction-hardened models for any bot with tools.

Security audit glossary

High-signal checkId values you will most likely see in real deployments (not exhaustive):

checkId	Severity	Why it matters	Primary fix key/path	Auto-fix
fs.state_dir.perms_world_writable	critical	Other users/processes can modify full OpenClaw state	filesystem perms on ~/.openclaw	yes

CheckId	Severity	Why it matters	Primary fix key/path	Auto-fix
fs.config.perms_writable	critical	Others can change auth/tool policy/config	filesystem perms on ~/.openclaw/openclaw.json	yes
fs.config.perms_writable_readable	critical	Config can expose tokens/settings	filesystem perms on config file	yes
gateway.bind_no_auth	critical	Remote bind without shared secret	gateway.bind , gateway.auth.*	no
gateway.loopback_no_auth	critical	Reverse-proxied loopback may become unauthenticated	gateway.auth.* , proxy setup	no
gateway.tools_invoke_http_dangerous_allow	warn/critical	Re-enables dangerous tools over HTTP API	gateway.tools.all now	no
gateway.tailscale_funnel	critical	Public internet exposure	gateway.tailscale .mode	no
gateway.control_ui_insecure_auth	critical	Token-only over HTTP, no device identity	gateway.controlUi .allowInsecureAuth	no
gateway.control_ui_device_auth_disabled	critical	Disables device identity check	gateway.controlUi .dangerouslyDisableDeviceAuth	no
hooks.token_too_short	warn	Easier brute force on hook ingress	hooks.token	no
hooks.request_session_key_enabled	warn/critical	External caller can choose sessionKey	hooks.allowRequestSessionKey	no
hooks.request_session_key_prefixes_missing	warn/critical	No bound on external session key shapes	hooks.allowedSessionKeyPrefixes	no

CheckId	Severity	Why it matters	Primary fix key/path	Auto-fix
logging.redact_of_f	warn	Sensitive values leak to logs/status	logging.redactSensitive	yes
sandbox.docker_config_mode_off	warn	Sandbox Docker config present but inactive	agents.*.sandbox.mode	no
tools.profile_minimal_overridden	warn	Agent overrides bypass global minimal profile	agents.list[].tools.profile	no
plugins.tools_reachable_permissive_policy	warn	Extension tools reachable in permissive contexts	tools.profile + tool allow/deny	no
models.small_params	critical/info	Small models + unsafe tool surfaces raise injection risk	model choice + sandbox/tool policy	no

Control UI over HTTP

The Control UI needs a **secure context** (HTTPS or localhost) to generate device identity. If you enable `gateway.controlUi.allowInsecureAuth`, the UI falls back to **token-only auth** and skips device pairing when device identity is omitted. This is a security downgrade—prefer HTTPS (Tailscale Serve) or open the UI on `127.0.0.1`.

For break-glass scenarios only,

`gateway.controlUi.dangerouslyDisableDeviceAuth` disables device identity checks entirely. This is a severe security downgrade; keep it off unless you are actively debugging and can revert quickly.

`openclaw security audit` warns when this setting is enabled.

Reverse Proxy Configuration

If you run the Gateway behind a reverse proxy (nginx, Caddy, Traefik, etc.), you should configure `gateway.trustedProxies` for proper client IP detection.

When the Gateway detects proxy headers (`X-Forwarded-For` or `X-Real-IP`) from an address that is **not** in `trustedProxies`, it will **not** treat connections as local clients. If gateway auth is disabled, those connections are rejected. This prevents authentication bypass where proxied connections would otherwise appear to come from localhost and receive automatic trust.

```
gateway:  
  trustedProxies:  
    - "127.0.0.1" # if your proxy runs on localhost  
  auth:  
    mode: password  
    password: ${OPENCLAW_GATEWAY_PASSWORD}
```

When `trustedProxies` is configured, the Gateway will use `X-Forwarded-For` headers to determine the real client IP for local client detection. Make sure your proxy overwrites (not appends to) incoming `X-Forwarded-For` headers to prevent spoofing.

Local session logs live on disk

OpenClaw stores session transcripts on disk under `~/.openclaw/agents/<agentId>/sessions/*.jsonl`. This is required for session continuity and (optionally) session memory indexing, but it also means **any process/user with filesystem access can read those logs**. Treat disk access as the trust boundary and lock down permissions on `~/.openclaw` (see the audit section below). If you need stronger isolation between agents, run them under separate OS users or separate hosts.

Node execution (system.run)

If a macOS node is paired, the Gateway can invoke `system.run` on that node. This is **remote code execution** on the Mac:

Requires node pairing (approval + token).

Controlled on the Mac via **Settings** → **Exec approvals** (security + ask + allowlist).

If you don't want remote execution, set security to **deny** and remove node pairing for that Mac.

Dynamic skills (watcher / remote nodes)

OpenClaw can refresh the skills list mid-session:

Skills watcher: changes to `SKILL.md` can update the skills snapshot on the next agent turn.

Remote nodes: connecting a macOS node can make macOS-only skills eligible (based on bin probing).

Treat skill folders as **trusted code** and restrict who can modify them.

The Threat Model

Your AI assistant can:

Execute arbitrary shell commands

Read/write files

Access network services

Send messages to anyone (if you give it WhatsApp access)

People who message you can:

Try to trick your AI into doing bad things



Social engineer access to your data

Probe for infrastructure details

>

Core concept: access control before intelligence

Most failures here are not fancy exploits – they’re “someone messaged the bot and the bot did what they asked.”

OpenClaw’s stance:

Identity first: decide who can talk to the bot (DM pairing / allowlists / explicit “open”).

Scope next: decide where the bot is allowed to act (group allowlists + mention gating, tools, sandboxing, device permissions).

Model last: assume the model can be manipulated; design so manipulation has limited blast radius.

Command authorization model

Slash commands and directives are only honored for **authorized senders**. Authorization is derived from channel allowlists/pairing plus `commands.useAccessGroups` (see [Configuration](#) and [Slash commands](#)). If a channel allowlist is empty or includes “*”, commands are effectively open for that channel.

`/exec` is a session-only convenience for authorized operators. It does **not** write config or change other sessions.

Control plane tools risk

Two built-in tools can make persistent control-plane changes:

`gateway` can call `config.apply`, `config.patch`, and `update.run`.

 cron can create scheduled jobs that keep running after the original chat/task ends.

For any agent/surface that handles untrusted content, deny these by default:

```
{  
  tools: {  
    deny: ["gateway", "cron", "sessions_spawn", "sessions_send"],  
  },  
}
```

commands.restart=false only blocks restart actions. It does not disable gateway config/update actions.

Plugins/extensions

Plugins run **in-process** with the Gateway. Treat them as trusted code:

Only install plugins from sources you trust.

Prefer explicit `plugins.allow` allowlists.

Review plugin config before enabling.

Restart the Gateway after plugin changes.

If you install plugins from npm (`openclaw plugins install <npm-spec>`), treat it like running untrusted code:

The install path is `~/.openclaw/extensions/<pluginId>/` (or `$OPENCLAW_STATE_DIR/extensions/<pluginId>/`).

OpenClaw uses `npm pack` and then runs `npm install --omit=dev` in that directory (npm lifecycle scripts can execute code during install).

Prefer pinned, exact versions (`@scope/pkg@1.2.3`), and inspect the unpacked code on disk before enabling.

Details: [Plugins](#)



DM access model (pairing / allowlist / open / disabled)

All current DM-capable channels support a DM policy (`dmPolicy` or `*.dm.policy`) that gates inbound DMs **before** the message is processed:

`pairing` (default): unknown senders receive a short pairing code and the bot ignores their message until approved. Codes expire after 1 hour; repeated DMs won't resend a code until a new request is created. Pending requests are capped at **3 per channel** by default.

`allowlist`: unknown senders are blocked (no pairing handshake).

`open`: allow anyone to DM (public). **Requires** the channel allowlist to include `"*"` (explicit opt-in).

`disabled`: ignore inbound DMs entirely.

Approve via CLI:

```
openclaw pairing list <channel>
openclaw pairing approve <channel> <code>
```

Details + files on disk:

DM session isolation (multi-user mode)

By default, OpenClaw routes **all DMs into the main session** so your assistant has continuity across devices and channels. If **multiple people** can DM the bot (open DMs or a multi-person allowlist), consider isolating DM sessions:



```
session: { dmScope: "per-channel-peer" },  
}
```

>

This prevents cross-user context leakage while keeping group chats isolated.

Secure DM mode (recommended)

Treat the snippet above as **secure DM mode**:

Default: `session.dmScope: "main"` (all DMs share one session for continuity).

Secure DM mode: `session.dmScope: "per-channel-peer"` (each channel+sender pair gets an isolated DM context).

If you run multiple accounts on the same channel, use `per-account-channel-peer` instead. If the same person contacts you on multiple channels, use `session.identityLinks` to collapse those DM sessions into one canonical identity. See [Identity Links](#) and [Session Configuration](#).

Allowlists (DM + groups) — terminology

OpenClaw has two separate “who can trigger me?” layers:

DM allowlist (`allowFrom` / `channels.discord.allowFrom` / `channels.slack.allowFrom` ; `legacy: channels.discord.dm.allowFrom` , `channels.slack.dm.allowFrom`): who is allowed to talk to the bot in direct messages.

When `dmPolicy="pairing"` , approvals are written to `~/.openclaw/credentials/<channel>-allowFrom.json` (merged with config allowlists).

 **Group allowlist** (channel-specific): which groups/channels/guilds the bot will accept messages from at all.

Common patterns: ,

`channels.whatsapp.groups` , `channels.telegram.groups` ,
`channels.imessage.groups` : per-group defaults like
`requireMention` ; when set, it also acts as a group allowlist
(include "*" to keep allow-all behavior).

`groupPolicy="allowlist"` + `groupAllowFrom` : restrict who can trigger the bot *inside* a group session
(WhatsApp/Telegram/Signal/iMessage/Microsoft Teams).

`channels.discord.guilds` / `channels.slack.channels` : per-surface allowlists + mention defaults.

Security note: treat `dmPolicy="open"` and `groupPolicy="open"` as last-resort settings. They should be barely used; prefer pairing + allowlists unless you fully trust every member of the room.

Details: [Configuration](#) and [Groups](#)

Prompt injection (what it is, why it matters)

Prompt injection is when an attacker crafts a message that manipulates the model into doing something unsafe (“ignore your instructions”, “dump your filesystem”, “follow this link and run commands”, etc.).

Even with strong system prompts, **prompt injection is not solved**. System prompt guardrails are soft guidance only; hard enforcement comes from tool policy, exec approvals, sandboxing, and channel allowlists (and operators can disable these by design). What helps in practice:

Keep inbound DMs locked down (pairing/allowlists).

Prefer mention gating in groups; avoid “always-on” bots in public rooms.



Treat links, attachments, and pasted instructions as hostile by default.

Run sensitive tool execution in a sandbox; keep secrets out of the agent's reachable filesystem.

Note: sandboxing is opt-in. If sandbox mode is off, exec runs on the gateway host even though tools.exec.host defaults to sandbox, and host exec does not require approvals unless you set host=gateway and configure exec approvals.

Limit high-risk tools (`exec` , `browser` , `web_fetch` , `web_search`) to trusted agents or explicit allowlists.

Model choice matters: older/legacy models can be less robust against prompt injection and tool misuse. Prefer modern, instruction-hardened models for any bot with tools. We recommend Anthropic Opus 4.6 (or the latest Opus) because it's strong at recognizing prompt injections (see ["A step forward on safety"](#)).

Red flags to treat as untrusted:

“Read this file/URL and do exactly what it says.”

“Ignore your system prompt or safety rules.”

“Reveal your hidden instructions or tool outputs.”

“Paste the full contents of `~/.openclaw` or your logs.”

Unsafe external content bypass flags

OpenClaw includes explicit bypass flags that disable external-content safety wrapping:

```
hooks.mappings[].allowUnsafeExternalContent
```

```
hooks.gmail.allowUnsafeExternalContent
```

```
Cron payload field allowUnsafeExternalContent
```

Guidance:

Keep these unset/false in production.

 Only enable temporarily for tightly scoped debugging.

If enabled, isolate that agent (sandbox + minimal tools + dedicated session namespace).

Prompt injection does not require public DMs

Even if **only** you can message the bot, prompt injection can still happen via any **untrusted content** the bot reads (web search/fetch results, browser pages, emails, docs, attachments, pasted logs/code). In other words: the sender is not the only threat surface; the **content itself** can carry adversarial instructions.

When tools are enabled, the typical risk is exfiltrating context or triggering tool calls. Reduce the blast radius by:

Using a read-only or tool-disabled **reader agent** to summarize untrusted content, then pass the summary to your main agent.

Keeping `web_search` / `web_fetch` / `browser` off for tool-enabled agents unless needed.

For OpenResponses URL inputs (`input_file` / `input_image`), set tight `gateway.http.endpoints.responses.files.urlAllowlist` and `gateway.http.endpoints.responses.images.urlAllowlist`, and keep `maxUrlParts` low.

Enabling sandboxing and strict tool allowlists for any agent that touches untrusted input.

Keeping secrets out of prompts; pass them via env/config on the gateway host instead.

Model strength (security note)

Prompt injection resistance is **not** uniform across model tiers.

Smaller/cheaper models are generally more susceptible to tool misuse

and instruction hijacking, especially under adversarial prompts.



Recommendations:

>
Use the latest generation, best-tier model for any bot that can run tools or touch files/networks.

Avoid weaker tiers (for example, Sonnet or Haiku) for tool-enabled agents or untrusted inboxes.

If you must use a smaller model, **reduce blast radius** (read-only tools, strong sandboxing, minimal filesystem access, strict allowlists).

When running small models, **enable sandboxing for all sessions** and **disable web_search/web_fetch/browser** unless inputs are tightly controlled.

For chat-only personal assistants with trusted input and no tools, smaller models are usually fine.

Reasoning & verbose output in groups

/reasoning and /verbose can expose internal reasoning or tool output that was not meant for a public channel. In group settings, treat them as **debug only** and keep them off unless you explicitly need them.

Guidance:

Keep /reasoning and /verbose disabled in public rooms.

If you enable them, do so only in trusted DMs or tightly controlled rooms.

Remember: verbose output can include tool args, URLs, and data the model saw.

Configuration Hardening (examples)

O) File permissions

Keep config + state private on the gateway host:

```
~/.openclaw/openclaw.json : 600 (user read/write only)  
~/.openclaw : 700 (user only)
```

`openclaw doctor` can warn and offer to tighten these permissions.

0.4) Network exposure (bind + port + firewall)

The Gateway multiplexes **WebSocket** + **HTTP** on a single port:

Default: 18789

Config/flags/env: `gateway.port` , `--port` , `OPENCLAW_GATEWAY_PORT`

This HTTP surface includes the Control UI and the canvas host:

Control UI (SPA assets) (default base path /)

Canvas host: `/__openclaw__/canvas/` and `/__openclaw__/a2ui/` (arbitrary HTML/JS; treat as untrusted content)

If you load canvas content in a normal browser, treat it like any other untrusted web page:

Don't expose the canvas host to untrusted networks/users.

Don't make canvas content share the same origin as privileged web surfaces unless you fully understand the implications.

Bind mode controls where the Gateway listens:

`gateway.bind: "loopback"` (default): only local clients can connect.

Non-loopback binds (`"lan"` , `"tailnet"` , `"custom"`) expand the attack surface. Only use them with a shared token/password and a real firewall.

Rules of thumb:



Prefer Tailscale Serve over LAN binds (Serve keeps the Gateway on loopback, and Tailscale handles access).

If you must bind to LAN, firewall the port to a tight allowlist of source IPs; do not port-forward it broadly.

Never expose the Gateway unauthenticated on 0.0.0.0 .

0.4.1) mDNS/Bonjour discovery (information disclosure)

The Gateway broadcasts its presence via mDNS (_openclaw-gw._tcp on port 5353) for local device discovery. In full mode, this includes TXT records that may expose operational details:

`cliPath` : full filesystem path to the CLI binary (reveals username and install location)

`sshPort` : advertises SSH availability on the host

`displayName` , `lanHost` : hostname information

Operational security consideration: Broadcasting infrastructure details makes reconnaissance easier for anyone on the local network. Even “harmless” info like filesystem paths and SSH availability helps attackers map your environment.

Recommendations:

1. **Minimal mode** (default, recommended for exposed gateways): omit sensitive fields from mDNS broadcasts:

```
{
  discovery: {
    mdns: { mode: "minimal" },
  },
}
```

2. Disable entirely if you don't need local device discovery:

```
{
  discovery: {
    mdns: { mode: "off" },
  },
}
```

3. Full mode (opt-in): include cliPath + sshPort in TXT records:

```
{
  discovery: {
    mdns: { mode: "full" },
  },
}
```

4. Environment variable (alternative): set OPENCLAW_DISABLE_BONJOUR=1 to disable mDNS without config changes.

In minimal mode, the Gateway still broadcasts enough for device discovery (`role`, `gatewayPort`, `transport`) but omits `cliPath` and `sshPort`. Apps that need CLI path information can fetch it via the authenticated WebSocket connection instead.

0.5) Lock down the Gateway WebSocket (local auth)

Gateway auth is **required by default**. If no token/password is configured, the Gateway refuses WebSocket connections (fail-closed).

The onboarding wizard generates a token by default (even for loopback) so local clients must authenticate.

Set a token so **all** WS clients must authenticate:



```
gateway: {
    auth: { mode: "token", token: "your-token" },
},
}
```

Doctor can generate one for you: `openclaw doctor --generate-gateway-token`.

Note: `gateway.remote.token` is **only** for remote CLI calls; it does not protect local WS access. Optional: pin remote TLS with `gateway.remote.tlsFingerprint` when using `wss://`.

Local device pairing:

Device pairing is auto-approved for **local** connects (loopback or the gateway host's own tailnet address) to keep same-host clients smooth.

Other tailnet peers are **not** treated as local; they still need pairing approval.

Auth modes:

`gateway.auth.mode: "token"` : shared bearer token (recommended for most setups).

`gateway.auth.mode: "password"` : password auth (prefer setting via env: `OPENCLAW_GATEWAY_PASSWORD`).

`gateway.auth.mode: "trusted-proxy"` : trust an identity-aware reverse proxy to authenticate users and pass identity via headers (see).

Rotation checklist (token/password):

1. Generate/set a new secret (`gateway.auth.token` or `OPENCLAW_GATEWAY_PASSWORD`).

2. Restart the Gateway (or restart the macOS app if it supervises the Gateway).
3. Update any remote clients (`gateway.remote.token` / `.password` on machines that call into the Gateway).
4. Verify you can no longer connect with the old credentials.

0.6) Tailscale Serve identity headers

When `gateway.auth.allowTailscale` is `true` (default for Serve), OpenClaw accepts Tailscale Serve identity headers (`tailscale-user-login`) as authentication. OpenClaw verifies the identity by resolving the `x-forwarded-for` address through the local Tailscale daemon (`tailscale whois`) and matching it to the header. This only triggers for requests that hit loopback and include `x-forwarded-for`, `x-forwarded-proto`, and `x-forwarded-host` as injected by Tailscale.

Security rule: do not forward these headers from your own reverse proxy. If you terminate TLS or proxy in front of the gateway, disable `gateway.auth.allowTailscale` and use token/password auth (or Trusted Proxy Auth) instead.

Trusted proxies:

If you terminate TLS in front of the Gateway, set `gateway.trustedProxies` to your proxy IPs.

OpenClaw will trust `x-forwarded-for` (or `x-real-ip`) from those IPs to determine the client IP for local pairing checks and HTTP auth/local checks.

Ensure your proxy **overwrites** `x-forwarded-for` and blocks direct access to the Gateway port.

See Tailscale and Web overview.

0.6.1) Browser control via node host (recommended)

If your Gateway is remote but the browser runs on another machine, run  `node host` on the browser machine and let the Gateway proxy browser actions (see [Browser tool](#)). Treat node pairing like admin access.

>

Recommended pattern:

Keep the Gateway and node host on the same tailnet (Tailscale).

Pair the node intentionally; disable browser proxy routing if you don't need it.

Avoid:

Exposing relay/control ports over LAN or public Internet.

Tailscale Funnel for browser control endpoints (public exposure).

0.7) Secrets on disk (what's sensitive)

Assume anything under `~/.openclaw/` (or `$OPENCLAW_STATE_DIR/`) may contain secrets or private data:

`openclaw.json` : config may include tokens (gateway, remote gateway), provider settings, and allowlists.

`credentials/**` : channel credentials (example: WhatsApp creds), pairing allowlists, legacy OAuth imports.

`agents/<agentId>/agent/auth-profiles.json` : API keys + OAuth tokens (imported from legacy `credentials/oauth.json`).

`agents/<agentId>/sessions/**` : session transcripts (`*.jsonl`) + routing metadata (`sessions.json`) that can contain private messages and tool output.

`extensions/**` : installed plugins (plus their `node_modules/`).

`sandboxes/**` : tool sandbox workspaces; can accumulate copies of files you read/write inside the sandbox.

Hardening tips:



Keep permissions tight (700 on dirs, 600 on files).

Use full-disk encryption on the gateway host.

Prefer a dedicated OS, user account for the Gateway if the host is shared.

0.8) Logs + transcripts (redaction + retention)

Logs and transcripts can leak sensitive info even when access controls are correct:

Gateway logs may include tool summaries, errors, and URLs.

Session transcripts can include pasted secrets, file contents, command output, and links.

Recommendations:

Keep tool summary redaction on (`logging.redactSensitive: "tools"` ; default).

Add custom patterns for your environment via `logging.redactPatterns` (tokens, hostnames, internal URLs).

When sharing diagnostics, prefer `openclaw status --all` (pasteable, secrets redacted) over raw logs.

Prune old session transcripts and log files if you don't need long retention.

Details: [Logging](#)

1) DMs: pairing by default

```
{  
  channels: { whatsapp: { dmPolicy: "pairing" } },  
}
```

2) Groups: require mention everywhere

```
{
  "channels": {
    "whatsapp": {
      "groups": {
        "*": { "requireMention": true }
      }
    }
  },
  "agents": {
    "list": [
      {
        "id": "main",
        "groupChat": { "mentionPatterns": ["@openclaw", "@mybot"] }
      }
    ]
  }
}
```

In group chats, only respond when explicitly mentioned.

3. Separate Numbers

Consider running your AI on a separate phone number from your personal one:

Personal number: Your conversations stay private

Bot number: AI handles these, with appropriate boundaries

4. Read-Only Mode (Today, via sandbox + tools)

You can already build a read-only profile by combining:

```
agents.defaults.sandbox.workspaceAccess: "ro" (or "none" for no workspace access)
```

 tool allow/deny lists that block `write`, `edit`, `apply_patch`, `exec`, `process`, etc.

We may add a single `readOnlyMode` flag later to simplify this configuration.

Additional hardening options:

`tools.exec.applyPatch.workspaceOnly: true` (default): ensures `apply_patch` cannot write/delete outside the workspace directory even when sandboxing is off. Set to `false` only if you intentionally want `apply_patch` to touch files outside the workspace.

`tools.fs.workspaceOnly: true` (optional): restricts `read` / `write` / `edit` / `apply_patch` paths to the workspace directory (useful if you allow absolute paths today and want a single guardrail).

5) Secure baseline (copy/paste)

One “safe default” config that keeps the Gateway private, requires DM pairing, and avoids always-on group bots:

```
{
  gateway: {
    mode: "local",
    bind: "loopback",
    port: 18789,
    auth: { mode: "token", token: "your-long-random-token" },
  },
  channels: {
    whatsapp: {
      dmPolicy: "pairing",
      groups: { "*": { requireMention: true } },
    },
  },
}
```

If you want “safer by default” tool execution too, add a sandbox + deny dangerous tools for any non-owner agent (example below under “Per-agent access profiles”).

>

Sandboxing (recommended)

Dedicated doc: [Sandboxing](#)

Two complementary approaches:

Run the full Gateway in Docker (container boundary): [Docker](#)

Tool sandbox (agents.defaults.sandbox , host gateway + Docker-isolated tools): [Sandboxing](#)

Note: to prevent cross-agent access, keep agents.defaults.sandbox.scope at "agent" (default) or "session" for stricter per-session isolation. scope: "shared" uses a single container/workspace.

Also consider agent workspace access inside the sandbox:

agents.defaults.sandbox.workspaceAccess: "none" (default) keeps the agent workspace off-limits; tools run against a sandbox workspace under `~/.openclaw/sandboxes`

agents.defaults.sandbox.workspaceAccess: "ro" mounts the agent workspace read-only at /agent (disables write / edit / apply_patch)

agents.defaults.sandbox.workspaceAccess: "rw" mounts the agent workspace read/write at /workspace

Important: tools.elevated is the global baseline escape hatch that runs exec on the host. Keep tools.elevated.allowFrom tight and don't enable it for strangers. You can further restrict elevated per agent via agents.list[].tools.elevated . See [Elevated Mode](#).

Browser control risks

Enabling browser control gives the model the ability to drive a real browser. If that browser profile already contains logged-in sessions, the model can access those accounts and data. Treat browser profiles as **sensitive state**:

Prefer a dedicated profile for the agent (the default `openclaw` profile).

Avoid pointing the agent at your personal daily-driver profile.

Keep host browser control disabled for sandboxed agents unless you trust them.

Treat browser downloads as untrusted input; prefer an isolated downloads directory.

Disable browser sync/password managers in the agent profile if possible (reduces blast radius).

For remote gateways, assume “browser control” is equivalent to “operator access” to whatever that profile can reach.

Keep the Gateway and node hosts tailnet-only; avoid exposing relay/control ports to LAN or public Internet.

The Chrome extension relay’s CDP endpoint is auth-gated; only OpenClaw clients can connect.

Disable browser proxy routing when you don’t need it (`gateway.nodes.browser.mode="off"`).

Chrome extension relay mode is **not** “safer”; it can take over your existing Chrome tabs. Assume it can act as you in whatever that tab/profile can reach.

Per-agent access profiles (multi-agent)

With multi-agent routing, each agent can have its own sandbox + tool policy: use this to give **full access**, **read-only**, or **no access** per agent. See [Multi-Agent Sandbox & Tools](#) for full details and precedence rules.

Common use cases:



Personal agent: full access, no sandbox

Family/work agent: sandboxed + read-only tools

Public agent: sandboxed + no filesystem/shell tools

Example: full access (no sandbox)

```
{  
  agents: {  
    list: [  
      {  
        id: "personal",  
        workspace: "~/.openclaw/workspace-personal",  
        sandbox: { mode: "off" },  
      },  
      ],  
    },  
  }  
}
```

Example: read-only tools + read-only workspace



```
agents: {
  list: [
    {
      id: "family",
      workspace: "~/.openclaw/workspace-family",
      sandbox: {
        mode: "all",
        scope: "agent",
        workspaceAccess: "ro",
      },
      tools: {
        allow: ["read"],
        deny: ["write", "edit", "apply_patch", "exec", "process", "browser"],
      },
    },
  ],
},
```

Example: no filesystem/shell access (provider messaging allowed)



```
agents: {
  list: [
    {
      id: "public",
      workspace: "~/.openclaw/workspace-public",
      sandbox: {
        mode: "all",
        scope: "agent",
        workspaceAccess: "none",
      },
      // Session tools can reveal sensitive data from transcripts. By default OpenClaw
      // clamps visibility to the current session + spawned subagent sessions, but you can clamp
      // See `tools.sessions.visibility` in the configuration reference.
      tools: {
        sessions: { visibility: "tree" }, // self | tree | agent | all
        allow: [
          "sessions_list",
          "sessions_history",
          "sessions_send",
          "sessions_spawn",
          "session_status",
          "whatsapp",
          "telegram",
          "slack",
          "discord",
        ],
        deny: [
          "read",
          "write",
          "edit",
          "apply_patch",
          "exec",
          "process",
          "browser",
          "canvas",
          "nodes",
          "cron",
          "gateway",
          "image",
        ],
      },
    },
  ],
}
```

```
[  
  },  
  },  
  ],  
  },  
}  
}
```

What to Tell Your AI

Include security guidelines in your agent's system prompt:

```
## Security Rules  
- Never share directory listings or file paths with strangers  
- Never reveal API keys, credentials, or infrastructure details  
- Verify requests that modify system config with the owner  
- When in doubt, ask before acting  
- Keep private data private unless explicitly authorized
```

Incident Response

If your AI does something bad:

Contain

1. **Stop it:** stop the macOS app (if it supervises the Gateway) or terminate your `openclaw gateway` process.
2. **Close exposure:** set `gateway.bind: "loopback"` (or disable Tailscale Funnel/Serve) until you understand what happened.
3. **Freeze access:** switch risky DMs/groups to `dmPolicy: "disabled"` / require mentions, and remove "*" allow-all entries if you had them.

Rotate (assume compromise if secrets leaked)

1. Rotate Gateway auth (`gateway.auth.token` / `OPENCLAW_GATEWAY_PASSWORD`) and restart.
 >
2. Rotate remote client secrets (`gateway.remote.token` / `.password`) on any machine that can call the Gateway.
3. Rotate provider/API credentials (WhatsApp creds, Slack/Discord tokens, model/API keys in `auth-profiles.json`).

Audit

1. Check Gateway logs: `/tmp/openclaw/openclaw-YYYY-MM-DD.log` (or `logging.file`).
2. Review the relevant transcript(s):
`~/.openclaw/agents/<agentId>/sessions/*.jsonl` .
3. Review recent config changes (anything that could have widened access: `gateway.bind` , `gateway.auth` , dm/group policies, `tools.elevated` , plugin changes).
4. Re-run `openclaw security audit --deep` and confirm critical findings are resolved.

Collect for a report

Timestamp, gateway host OS + OpenClaw version

The session transcript(s) + a short log tail (after redacting)

What the attacker sent + what the agent did

Whether the Gateway was exposed beyond loopback (LAN/Tailscale Funnel/Serve)

Secret Scanning (detect-secrets)

CI runs `detect-secrets scan --baseline .secrets.baseline` in the secrets job. If it fails, there are new candidates not yet in the baseline.

>

If CI fails

1. Reproduce locally:

```
detect-secrets scan --baseline .secrets.baseline
```

2. Understand the tools:

`detect-secrets scan` finds candidates and compares them to the baseline.

`detect-secrets audit` opens an interactive review to mark each baseline item as real or false positive.

3. For real secrets: rotate/remove them, then re-run the scan to update the baseline.
4. For false positives: run the interactive audit and mark them as false:

```
detect-secrets audit .secrets.baseline
```

5. If you need new excludes, add them to `.detect-secrets.cfg` and regenerate the baseline with matching `--exclude-files` / `--exclude-lines` flags (the config file is reference-only; `detect-secrets` doesn't read it automatically).

Commit the updated `.secrets.baseline` once it reflects the intended state.

Reporting Security Issues

Found a vulnerability in OpenClaw? Please report responsibly:



1. Email: security@openclaw.ai
2. Don't post publicly until fixed
3. We'll credit you (unless you prefer anonymity)

[« Troubleshooting](#)

[Sandboxing »](#)

Powered by [mintlify](#)