



☰ Compaction internals > **Session Management Deep Dive**

[Compaction internals](#)

Session Management Deep Dive

This document explains how OpenClaw manages sessions end-to-end:

Session routing (how inbound messages map to a `sessionKey`)

Session store (`sessions.json`) and what it tracks

Transcript persistence (`*.jsonl`) and its structure

Transcript hygiene (provider-specific fixups before runs)

Context limits (context window vs tracked tokens)

Compaction (manual + auto-compaction) and where to hook pre-compaction work

Silent housekeeping (e.g. memory writes that shouldn't produce user-visible output)

If you want a higher-level overview first, start with:

[/concepts/session](#)

[/concepts/compaction](#)

[/concepts/session-pruning](#)

[/reference/transcript-hygiene](#)

Source of truth: the Gateway

OpenClaw is designed around a single **Gateway process** that owns session state.

UIs (macOS app, web Control UI, TUI) should query the Gateway for session lists and token counts.

In remote mode, session files are on the remote host; “checking your local Mac files” won’t reflect what the Gateway is using.

Two persistence layers

OpenClaw persists sessions in two layers:

1. Session store (sessions.json)

Key/value map: `sessionKey -> SessionEntry`

Small, mutable, safe to edit (or delete entries)

Tracks session metadata (current session id, last activity, toggles, token counters, etc.)

2. Transcript (<sessionId>.jsonl)

Append-only transcript with tree structure (entries have `id` + `parentId`)

Stores the actual conversation + tool calls + compaction summaries

Used to rebuild the model context for future turns

On-disk locations

Per agent, on the Gateway host:



Store: `~/.openclaw/agents/<agentId>/sessions/sessions.json`

Transcripts: `~/.openclaw/agents/<agentId>/sessions/<sessionId>.jsonl`

Telegram topic sessions: `.../<sessionId>-topic-<threadId>.jsonl`

OpenClaw resolves these via `src/config/sessions.ts`.

Session keys (`sessionKey`)

A `sessionKey` identifies *which conversation bucket* you're in (routing + isolation).

Common patterns:

Main/direct chat (per agent): `agent:<agentId>:<mainKey>` (default `main`)

Group: `agent:<agentId>:<channel>:group:<id>`

Room/channel (Discord/Slack): `agent:<agentId>:<channel>:channel:<id>` or
`...:room:<id>`

Cron: `cron:<job.id>`

Webhook: `hook:<uuid>` (unless overridden)

The canonical rules are documented at [/concepts/session](#).

Session ids (`sessionId`)

Each `sessionKey` points at a current `sessionId` (the transcript file that continues the conversation).

Rules of thumb:



Reset (/new , /reset) creates a new sessionId for that sessionKey .

Daily reset (default 4:00 AM local time on the gateway host) creates a new sessionId on the next message after the reset boundary.

Idle expiry (session.reset.idleMinutes or legacy session.idleMinutes) creates a new sessionId when a message arrives after the idle window. When daily + idle are both configured, whichever expires first wins.

Implementation detail: the decision happens in initSessionState() in src/auto-reply/reply/session.ts .

Session store schema (sessions.json)

The store's value type is SessionEntry in src/config/sessions.ts .

Key fields (not exhaustive):

sessionId : current transcript id (filename is derived from this unless sessionFile is set)

updatedAt : last activity timestamp

sessionFile : optional explicit transcript path override

chatType : direct | group | room (helps UIs and send policy)

provider , subject , room , space , displayName : metadata for group/channel labeling

Toggles:

thinkingLevel , verboseLevel , reasoningLevel , elevatedLevel
sendPolicy (per-session override)

Model selection:

providerOverride , modelOverride , authProfileOverride



Token counters (best-effort / provider-dependent):

`inputTokens , outputTokens , totalTokens , contextTokens`

`compactionCount` : how often auto-compaction completed for this session key

`memoryFlushAt` : timestamp for the last pre-compaction memory flush

`memoryFlushCompactionCount` : compaction count when the last flush ran

The store is safe to edit, but the Gateway is the authority: it may rewrite or rehydrate entries as sessions run.

Transcript structure (*.jsonl)

Transcripts are managed by `@mariozechner/pi-coding-agent`'s `SessionManager`.

The file is JSONL:

First line: session header (`type: "session"` , includes `id` , `cwd` , `timestamp` , optional `parentSession`)

Then: session entries with `id` + `parentId` (tree)

Notable entry types:

`message` : user/assistant/toolResult messages

`custom_message` : extension-injected messages that *do* enter model context (can be hidden from UI)

`custom` : extension state that does *not* enter model context

`compaction` : persisted compaction summary with `firstKeptEntryId` and `tokensBefore`

`branch_summary` : persisted summary when navigating a tree branch

OpenClaw intentionally does **not** “fix up” transcripts; the Gateway uses `SessionManager` to read/write them.

>

Context windows vs tracked tokens

Two different concepts matter:

1. **Model context window**: hard cap per model (tokens visible to the model)
2. **Session store counters**: rolling stats written into `sessions.json` (used for /status and dashboards)

If you’re tuning limits:

The context window comes from the model catalog (and can be overridden via config).

`contextTokens` in the store is a runtime estimate/reporting value; don’t treat it as a strict guarantee.

For more, see [/token-use](#).

Compaction: what it is

Compaction summarizes older conversation into a persisted `compaction` entry in the transcript and keeps recent messages intact.

After compaction, future turns see:

The compaction summary

Messages after `firstKeptEntryId`

Compaction is **persistent** (unlike session pruning). See
 [concepts/session-pruning](#).

>

When auto-compaction happens (Pi runtime)

In the embedded Pi agent, auto-compaction triggers in two cases:

1. **Overflow recovery**: the model returns a context overflow error → compact → retry.
2. **Threshold maintenance**: after a successful turn, when:

`contextTokens > contextWindow - reserveTokens`

Where:

`contextWindow` is the model's context window

`reserveTokens` is headroom reserved for prompts + the next model output

These are Pi runtime semantics (OpenClaw consumes the events, but Pi decides when to compact).

Compaction settings (`reserveTokens` , `keepRecentTokens`)

Pi's compaction settings live in Pi settings:



```
compaction: {  
    enabled: true,  
    reserveTokens: 16384,  
    keepRecentTokens: 20000,  
},  
}
```

OpenClaw also enforces a safety floor for embedded runs:

If `compaction.reserveTokens < reserveTokensFloor`, OpenClaw bumps it.

Default floor is 20000 tokens.

Set `agents.defaults.compaction.reserveTokensFloor: 0` to disable the floor.

If it's already higher, OpenClaw leaves it alone.

Why: leave enough headroom for multi-turn “housekeeping” (like memory writes) before compaction becomes unavoidable.

Implementation: `ensurePiCompactionReserveTokens()` in `src/agents/pi-settings.ts` (called from `src/agents/pi-embedded-runner.ts`).

User-visible surfaces

You can observe compaction and session state via:

`/status` (in any chat session)

`openclaw status` (CLI)

`openclaw sessions / sessions --json`

Verbose mode: Auto-compaction complete + compaction count

Silent housekeeping (NO_REPLY)

OpenClaw supports “silent” turns for background tasks where the user should not see intermediate output.

Convention:

The assistant starts its output with `NO_REPLY` to indicate “do not deliver a reply to the user”.

OpenClaw strips/suppresses this in the delivery layer.

As of `2026.1.10`, OpenClaw also suppresses `draft/typing streaming` when a partial chunk begins with `NO_REPLY`, so silent operations don’t leak partial output mid-turn.

Pre-compaction “memory flush” (implemented)

Goal: before auto-compaction happens, run a silent agentic turn that writes durable state to disk (e.g. `memory/YYYY-MM-DD.md` in the agent workspace) so compaction can’t erase critical context.

OpenClaw uses the `pre-threshold flush` approach:

1. Monitor session context usage.
2. When it crosses a “soft threshold” (below Pi’s compaction threshold), run a silent “write memory now” directive to the agent.
3. Use `NO_REPLY` so the user sees nothing.

Config (`agents.defaults.compaction.memoryFlush`):

```
enabled (default: true)  
softThresholdTokens (default: 4000)
```



`prompt` (user message for the flush turn)

`systemPrompt` (extra system prompt appended for the flush turn)

Notes:

>

The default prompt/system prompt include a `NO_REPLY` hint to suppress delivery.

The flush runs once per compaction cycle (tracked in `sessions.json`).

The flush runs only for embedded Pi sessions (CLI backends skip it).

The flush is skipped when the session workspace is read-only (`workspaceAccess: "ro"` or `"none"`).

See [**Memory**](#) for the workspace file layout and write patterns.

Pi also exposes a `session_before_compact` hook in the extension API, but OpenClaw's flush logic lives on the Gateway side today.

Troubleshooting checklist

Session key wrong? Start with [**/concepts/session**](#) and confirm the `sessionKey` in `/status` .

Store vs transcript mismatch? Confirm the Gateway host and the store path from `openclaw status` .

Compaction spam? Check:

model context window (too small)

compaction settings (`reserveTokens` too high for the model window can cause earlier compaction)

tool-result bloat: enable/tune session pruning



Silent turns leaking? Confirm the reply starts with `NO_REPLY` (exact token) and you're on a build that includes the streaming suppression fix.

>

[◀ Node.js](#)[Setup ▶](#)

Powered by [mintlify](#)