



Hosting and deployment

Fly.io

Deploy OpenClaw on Fly.io

Goal: OpenClaw Gateway running on a Fly.io machine with persistent storage, automatic HTTPS, and Discord/channel access.

) What you need

flyctl CLI installed

Fly.io account (free tier works)

Model auth: Anthropic API key (or other provider keys)

Channel credentials: Discord bot token, Telegram token, etc.

Beginner quick path

1. Clone repo → customize `fly.toml`
2. Create app + volume → set secrets
3. Deploy with `fly deploy`
4. SSH in to create config or use Control UI

1) Create the Fly app

```
# Clone the repo
git clone https://github.com/openclaw/openclaw.git
cd openclaw
>

# Create a new Fly app (pick your own name)
fly apps create my-openclaw

# Create a persistent volume (1GB is usually enough)
fly volumes create openclaw_data --size 1 --region iad
```

Tip: Choose a region close to you. Common options: `lhr` (London), `iad` (Virginia), `sjc` (San Jose).

2) Configure `fly.toml`

Edit `fly.toml` to match your app name and requirements.

Security note: The default config exposes a public URL. For a hardened deployment with no public IP, see [this guide](#) or use `fly.private.toml`.

```

app = "my-openclaw" # Your app name
primary_region = "iad"

[build]
>
dockerfile = "Dockerfile"

[env]
NODE_ENV = "production"
OPENCLAW_PREFER_PNPM = "1"
OPENCLAW_STATE_DIR = "/data"
NODE_OPTIONS = "--max-old-space-size=1536"

[processes]
app = "node dist/index.js gateway --allow-unconfigured --port 3000 --bind lan"

[http_service]
internal_port = 3000
force_https = true
auto_stop_machines = false
auto_start_machines = true
min_machines_running = 1
processes = ["app"]

[[vm]]
size = "shared-cpu-2x"
memory = "2048mb"

[mounts]
source = "openclaw_data"
destination = "/data"

```

Key settings:

Setting	Why
--bind lan	Binds to 0.0.0.0 so Fly's proxy can reach the gateway
--allow-unconfigured	Starts without a config file (you'll create one after)

Setting	Why
 Internal_port = 3000	Must match --port 3000 (or OPENCLAW_GATEWAY_PORT) for Fly health checks
memory = "2048mb"	512MB is too small; 2GB recommended
OPENCLAW_STATE_DIR = "/data"	Persists state on the volume

3) Set secrets

```
# Required: Gateway token (for non-loopback binding)
fly secrets set OPENCLAW_GATEWAY_TOKEN=$(openssl rand -hex 32)

# Model provider API keys
fly secrets set ANTHROPIC_API_KEY=sk-ant-...

# Optional: Other providers
fly secrets set OPENAI_API_KEY=sk-...
fly secrets set GOOGLE_API_KEY=...

# Channel tokens
fly secrets set DISCORD_BOT_TOKEN=MTQ...
```

Notes:

Non-loopback binds (--bind lan) require OPENCLAW_GATEWAY_TOKEN for security.

Treat these tokens like passwords.

Prefer env vars over config file for all API keys and tokens. This keeps secrets out of openclaw.json where they could be accidentally exposed or logged.

4) Deploy

 `fly deploy`

First deploy builds the Docker image (~2-3 minutes). Subsequent deploys are faster.

After deployment, verify:

`fly status`
`fly logs`

You should see:

```
[gateway] listening on ws://0.0.0.0:3000 (PID xxx)
[discord] logged in to discord as xxx
```

5) Create config file

SSH into the machine to create a proper config:

`fly ssh console`

Create the config directory and file:

```
mkdir -p /data
cat > /data/openclaw.json << 'EOF'
{
  "agents": {
    "defaults": {
      "model": {
        "primary": "anthropic/clause-opus-4-6",
        "fallbacks": ["anthropic/clause-sonnet-4-5", "openai/gpt-4o"]
      },
      "maxConcurrent": 4
    },
    "list": [
      {
        "id": "main",
        "default": true
      }
    ]
  },
  "auth": {
    "profiles": {
      "anthropic:default": { "mode": "token", "provider": "anthropic" },
      "openai:default": { "mode": "token", "provider": "openai" }
    }
  },
  "bindings": [
    {
      "agentId": "main",
      "match": { "channel": "discord" }
    }
  ],
  "channels": {
    "discord": {
      "enabled": true,
      "groupPolicy": "allowlist",
      "guilds": {
        "YOUR_GUILD_ID": {
          "channels": { "general": { "allow": true } },
          "requireMention": false
        }
      }
    }
  }
}
```

```
    },
  },
  "gateway": {
    "mode": "local",
    "bind": "auto"
  },
  "meta": {
    "lastTouchedVersion": "2026.1.29"
  }
}
EOF
```

Note: With `OPENCLAW_STATE_DIR=/data` , the config path is `/data/openclaw.json` .

Note: The Discord token can come from either:

Environment variable: `DISCORD_BOT_TOKEN` (recommended for secrets)

Config file: `channels.discord.token`

If using env var, no need to add token to config. The gateway reads `DISCORD_BOT_TOKEN` automatically.

Restart to apply:

```
exit
fly machine restart <machine-id>
```

6) Access the Gateway

Control UI

Open in browser:

```
fly open
```

Or visit <https://my-openclaw.fly.dev/>



Paste your gateway token (the one from `OPENCLAW_GATEWAY_TOKEN`) to authenticate.

>

Logs

```
fly logs           # Live logs  
fly logs --no-tail # Recent logs
```

SSH Console

```
fly ssh console
```

Troubleshooting

"App is not listening on expected address"

The gateway is binding to `127.0.0.1` instead of `0.0.0.0`.

Fix: Add `--bind lan` to your process command in `fly.toml`.

Health checks failing / connection refused

Fly can't reach the gateway on the configured port.

Fix: Ensure `internal_port` matches the gateway port (set `--port 3000` or `OPENCLAW_GATEWAY_PORT=3000`).

OOM / Memory Issues

Container keeps restarting or getting killed. Signs: `SIGABRT`, `v8::internal::Runtime_AllocateInYoungGeneration`, or silent restarts.

Fix: Increase memory in `fly.toml`:



```
[ [vm] ]
memory = "2048mb"      >
```

Or update an existing machine:

```
fly machine update <machine-id> --vm-memory 2048 -y
```

Note: 512MB is too small. 1GB may work but can OOM under load or with verbose logging. **2GB is recommended.**

Gateway Lock Issues

Gateway refuses to start with “already running” errors.

This happens when the container restarts but the PID lock file persists on the volume.

Fix: Delete the lock file:

```
fly ssh console --command "rm -f /data/gateway.*.lock"
fly machine restart <machine-id>
```

The lock file is at `/data/gateway.*.lock` (not in a subdirectory).

Config Not Being Read

If using `--allow-unconfigured`, the gateway creates a minimal config. Your custom config at `/data/openclaw.json` should be read on restart.

Verify the config exists:

 fly ssh console --command "cat /data/openclaw.json"

Writing Config via SSH

The `fly ssh console -C` command doesn't support shell redirection. To write a config file:

```
# Use echo + tee (pipe from local to remote)
echo '{"your":"config"}' | fly ssh console -C "tee /data/openclaw.json"

# Or use sftp
fly sftp shell
> put /local/path/config.json /data/openclaw.json
```

Note: `fly sftp` may fail if the file already exists. Delete first:

```
fly ssh console --command "rm /data/openclaw.json"
```

State Not Persisting

If you lose credentials or sessions after a restart, the state dir is writing to the container filesystem.

Fix: Ensure `OPENCLAW_STATE_DIR=/data` is set in `fly.toml` and redeploy.

Updates



```
# Pull latest changes  
git pull
```

```
# Redeploy >  
fly deploy
```

```
# Check health  
fly status  
fly logs
```

Updating Machine Command

If you need to change the startup command without a full redeploy:

```
# Get machine ID  
fly machines list  
  
# Update command  
fly machine update <machine-id> --command "node dist/index.js gateway --port 3000  
  
# Or with memory increase  
fly machine update <machine-id> --vm-memory 2048 --command "node dist/index.js ga
```

Note: After `fly deploy`, the machine command may reset to what's in `fly.toml`. If you made manual changes, re-apply them after deploy.

Private Deployment (Hardened)

By default, Fly allocates public IPs, making your gateway accessible at `https://your-app.fly.dev`. This is convenient but means your deployment is discoverable by internet scanners (Shodan, Censys, etc.).

For a hardened deployment with **no public exposure**, use the private template.

>

When to use private deployment

You only make **outbound** calls/messages (no inbound webhooks)

You use **ngrok or Tailscale** tunnels for any webhook callbacks

You access the gateway via **SSH, proxy, or WireGuard** instead of browser

You want the deployment **hidden from internet scanners**

Setup

Use `fly.private.toml` instead of the standard config:

```
# Deploy with private config
fly deploy -c fly.private.toml
```

Or convert an existing deployment:

```
# List current IPs
fly ips list -a my-openclaw

# Release public IPs
fly ips release <public-ipv4> -a my-openclaw
fly ips release <public-ipv6> -a my-openclaw

# Switch to private config so future deploys don't re-allocate public IPs
# (remove [http_service] or deploy with the private template)
fly deploy -c fly.private.toml

# Allocate private-only IPv6
fly ips allocate-v6 --private -a my-openclaw
```

After this, `fly ips list` should show only a private type IP:



VERSION	IP	TYPE	REGION
v6	fdaa:x:x:x:x::x	private	global

Accessing a private deployment

Since there's no public URL, use one of these methods:

Option 1: Local proxy (simplest)

```
# Forward local port 3000 to the app
fly proxy 3000:3000 -a my-openclaw

# Then open http://localhost:3000 in browser
```

Option 2: WireGuard VPN

```
# Create WireGuard config (one-time)
fly wireguard create

# Import to WireGuard client, then access via internal IPv6
# Example: http://[fdaa:x:x:x:x::x]:3000
```

Option 3: SSH only

```
fly ssh console -a my-openclaw
```

Webhooks with private deployment

If you need webhook callbacks (Twilio, Telnyx, etc.) without public exposure:

-  1. **ngrok tunnel** - Run ngrok inside the container or as a sidecar
2. **Tailscale Funnel** - Expose specific paths via Tailscale
3. **Outbound-only** - Some providers (Twilio) work fine for outbound calls without webhooks

Example voice-call config with ngrok:

```
{
  "plugins": {
    "entries": {
      "voice-call": {
        "enabled": true,
        "config": {
          "provider": "twilio",
          "tunnel": { "provider": "ngrok" },
          "webhookSecurity": {
            "allowedHosts": ["example.ngrok.app"]
          }
        }
      }
    }
  }
}
```

The ngrok tunnel runs inside the container and provides a public webhook URL without exposing the Fly app itself. Set `webhookSecurity.allowedHosts` to the public tunnel hostname so forwarded host headers are accepted.

Security benefits

Aspect	Public	Private
Internet scanners	Discoverable	Hidden
Direct attacks	Possible	Blocked

Aspect	Public	Private
 Control UI access	Browser	Proxy/VPN
Webhook delivery	> Direct	Via tunnel

Notes

Fly.io uses **x86 architecture** (not ARM)

The Dockerfile is compatible with both architectures

For WhatsApp/Telegram onboarding, use `fly ssh console`

Persistent data lives on the volume at `/data`

Signal requires Java + signal-cli; use a custom image and keep memory at 2GB+.

Cost

With the recommended config (`shared-cpu-2x` , 2GB RAM) :

~\$10-15/month depending on usage

Free tier includes some allowance

See [Fly.io pricing](#) for details.

[< Uninstall](#)

[Hetzner >](#)

Powered by [mintlify](#)