Providers

MiniMax

MiniMax is an AI company that builds the **M2/M2.1** model family. The current coding-focused release is **MiniMax M2.1** (December 23, 2025), built for real-world complex tasks.

Source: **MiniMax M2.1 release note**

Model overview (M2.1)

MiniMax highlights these improvements in M2.1:

Stronger **multi-language coding** (Rust, Java, Go, C++, Kotlin, Objective-C, TS/JS).

Better **web/app development** and aesthetic output quality (including native mobile).

Improved **composite instruction** handling for office-style workflows, building on interleaved thinking and integrated constraint execution.

More concise responses with lower token usage and faster iteration loops.

Stronger **tool/agent framework** compatibility and context management (Claude Code, Droid/Factory AI, Cline, Kilo Code, Roo Code, BlackBox).

Higher-quality **dialogue and technical writing** outputs.

MiniMax M2.1 vs MiniMax M2.1 Lightning

- **Speed:** Lightning is the “fast” variant in MiniMax’s pricing docs.
- **Cost:** Pricing shows the same input cost, but Lightning has higher output cost.
- **Coding plan routing:** The Lightning back-end isn’t directly available on the MiniMax coding plan. MiniMax auto-routes most requests to Lightning, but falls back to the regular M2.1 back-end during traffic spikes.

Choose a setup

MiniMax OAuth (Coding Plan) — recommended

Best for: quick setup with MiniMax Coding Plan via OAuth, no API key required.

Enable the bundled OAuth plugin and authenticate:

```
openclaw plugins enable minimax-portal-auth  # skip if already loaded.
openclaw gateway restart  # restart if gateway is already running
openclaw onboard --auth-choice minimax-portal
```

You will be prompted to select an endpoint:

- **Global** - International users (`api.minimax.io`)
- **CN** - Users in China (`api.minimaxi.com`)

See for details.

MiniMax M2.1 (API key)

Best for: hosted MiniMax with Anthropic-compatible API.

Configure via CLI:

Run `openclaw configure`

Select **Model/auth**

Choose **MiniMax M2.1**

```
{
  env: { MINIMAX_API_KEY: "sk-..." },
  agents: { defaults: { model: { primary: "minimax/MiniMax-M2.1" } } },
  models: {
    mode: "merge",
    providers: {
      minimax: {
        baseUrl: "https://api.minimax.io/anthropic",
        apiKey: "${MINIMAX_API_KEY}",
        api: "anthropic-messages",
        models: [
          {
            id: "MiniMax-M2.1",
            name: "MiniMax M2.1",
            reasoning: false,
            input: ["text"],
            cost: { input: 15, output: 60, cacheRead: 2, cacheWrite: 10 },
            contextWindow: 200000,
            maxTokens: 8192,
          },
        ],
      },
    },
  },
}
```

MiniMax M2.1 as fallback (Opus primary)

Best for: keep Opus 4.6 as primary, fail over to MiniMax M2.1.

```
{
  env: { MINIMAX_API_KEY: "sk-..." },
  agents: {
    defaults: {
      models: {
        "anthropic/claude-opus-4-6": { alias: "opus" },
        "minimax/MiniMax-M2.1": { alias: "minimax" },
      },
      model: {
        primary: "anthropic/claude-opus-4-6",
        fallbacks: ["minimax/MiniMax-M2.1"],
      },
    },
  },
}
```

Optional: Local via LM Studio (manual)

Best for: local inference with LM Studio. We have seen strong results with MiniMax M2.1 on powerful hardware (e.g. a desktop/server) using LM Studio's local server.

Configure manually via `openclaw.json` :

```
{
  agents: {
    defaults: {
      model: { primary: "lmstudio/minimax-m2.1-gs32" },
      models: { "lmstudio/minimax-m2.1-gs32": { alias: "Minimax" } },
    },
  },
  models: {
    mode: "merge",
    providers: {
      lmstudio: {
        baseUrl: "http://127.0.0.1:1234/v1",
        apiKey: "lmstudio",
        api: "openai-responses",
        models: [
          {
            id: "minimax-m2.1-gs32",
            name: "MiniMax M2.1 GS32",
            reasoning: false,
            input: ["text"],
            cost: { input: 0, output: 0, cacheRead: 0, cacheWrite: 0 },
            contextWindow: 196608,
            maxTokens: 8192,
          },
        ],
      },
    },
  },
}
```

Configure via `openclaw configure`

Use the interactive config wizard to set MiniMax without editing JSON:

1. Run `openclaw configure`.
2. Select **Model/auth**.
3. Choose **MiniMax M2.1**.

4. Pick your default model when prompted.

Configuration options

`models.providers.minimax.baseUrl` : prefer `https://api.minimax.io/anthropic` (Anthropic-compatible); `https://api.minimax.io/v1` is optional for OpenAI-compatible payloads.

`models.providers.minimax.api` : prefer `anthropic-messages` ; `openai-completions` is optional for OpenAI-compatible payloads.

`models.providers.minimax.apiKey` : MiniMax API key (`MINIMAX_API_KEY`).

`models.providers.minimax.models` : define `id` , `name` , `reasoning` , `contextWindow` , `maxTokens` , `cost` .

`agents.defaults.models` : alias models you want in the allowlist.

`models.mode` : keep `merge` if you want to add MiniMax alongside built-ins.

Notes

Model refs are `minimax/<model>` .

Coding Plan usage API: `https://api.minimaxi.com/v1/api/openplatform/coding_plan/remains` (requires a coding plan key).

Update pricing values in `models.json` if you need exact cost tracking.

Referral link for MiniMax Coding Plan (10% off): https://platform.minimax.io/subscribe/coding-plan?code=DbXJTRClnb&source=link

See /concepts/model-providers for provider rules.

Use `openclaw models list` and `openclaw models set minimax/MiniMax-M2.1` to switch.

Troubleshooting

“Unknown model: minimax/MiniMax-M2.1”

This usually means the **MiniMax provider isn’t configured** (no provider entry and no MiniMax auth profile/env key found). A fix for this detection is in **2026.1.12** (unreleased at the time of writing). Fix by:

- Upgrading to **2026.1.12** (or run from source `main`), then restarting the gateway.
- Running `openclaw configure` and selecting **MiniMax M2.1**, or
- Adding the `models.providers.minimax` block manually, or
- Setting `MINIMAX_API_KEY` (or a MiniMax auth profile) so the provider can be injected.

Make sure the model id is **case-sensitive**:

- `minimax/MiniMax-M2.1`
- `minimax/MiniMax-M2.1-lightning`

Then recheck with:

```
openclaw models list
```

‹ Moonshot AI OpenCode Zen ›