



## ☰ Configuration > Model Providers

### Configuration

## Model Providers

This page covers **LLM/model providers** (not chat channels like WhatsApp/Telegram). For model selection rules, see [/concepts/models](#).

## Quick rules

Model refs use `provider/model` (example: `opencode/clause-opus-4-6`).

If you set `agents.defaults.models`, it becomes the allowlist.

CLI helpers: `openclaw onboard` , `openclaw models list` , `openclaw models set <provider/model>` .

## API key rotation

Supports generic provider rotation for selected providers.

Configure multiple keys via:

`OPENCLAW_LIVE_<PROVIDER>_KEY` (single live override, highest priority)

`<PROVIDER>_API_KEYS` (comma or semicolon list)

`<PROVIDER>_API_KEY` (primary key)

`<PROVIDER>_API_KEY_*` (numbered list, e.g. `<PROVIDER>_API_KEY_1`)

For Google providers, `GOOGLE_API_KEY` is also included as fallback.

Key selection order preserves priority and deduplicates values.



Requests are retried with the next key only on rate-limit responses  
 (for example 429 , rate\_limit , quota , resource exhausted ).

Non-rate-limit failures fail immediately; no key rotation is  
 attempted.

When all candidate keys fail, the final error is returned from the  
 last attempt.

## Built-in providers (pi-ai catalog)

OpenClaw ships with the pi-ai catalog. These providers require no  
`models.providers` config; just set auth + pick a model.

### OpenAI

Provider: openai

Auth: OPENAI\_API\_KEY

Optional rotation: OPENAI\_API\_KEYS , OPENAI\_API\_KEY\_1 ,  
 OPENAI\_API\_KEY\_2 , plus OPENCLAW\_LIVE\_OPENAI\_KEY (single override)

Example model: openai/gpt-5.1-codex

CLI: `openclaw onboard --auth-choice openai-api-key`

```
{
  agents: { defaults: { model: { primary: "openai/gpt-5.1-codex" } } },
}
```

### Anthropic

Provider: anthropic

Auth: ANTHROPIC\_API\_KEY or claude setup-token

Optional rotation: ANTHROPIC\_API\_KEYS , ANTHROPIC\_API\_KEY\_1 ,  
 ANTHROPIC\_API\_KEY\_2 , plus OPENCLAW\_LIVE\_ANTHROPIC\_KEY (single override)



Example model: `anthropic/clause-opus-4-6`

CLI: `openclaw onboard --auth-choice token` (paste setup-token) or  
`openclaw models auth paste-token --provider anthropic`

```
{
  agents: { defaults: { model: { primary: "anthropic/clause-opus-4-6" } } },
}
```

## OpenAI Code (Codex)

Provider: `openai-codex`

Auth: OAuth (ChatGPT)

Example model: `openai-codex/gpt-5.3-codex`

CLI: `openclaw onboard --auth-choice openai-codex` or `openclaw models auth login --provider openai-codex`

```
{
  agents: { defaults: { model: { primary: "openai-codex/gpt-5.3-codex" } } },
}
```

## OpenCode Zen

Provider: `opencode`

Auth: `OPENCODE_API_KEY` (or `OPENCODE_ZEN_API_KEY`)

Example model: `opencode/clause-opus-4-6`

CLI: `openclaw onboard --auth-choice opencode-zen`

```
{
  agents: { defaults: { model: { primary: "opencode/clause-opus-4-6" } } },
}
```

## Google Gemini (API key)



Provider: google

Auth: GEMINI\_API\_KEY

Optional rotation: GEMINI\_API\_KEYS , GEMINI\_API\_KEY\_1 ,  
GEMINI\_API\_KEY\_2 , GOOGLE\_API\_KEY fallback, and  
OPENCLAW\_LIVE\_GEMINI\_KEY (single override)

Example model: google/gemini-3-pro-preview

CLI: openclaw onboard --auth-choice gemini-api-key

## Google Vertex, Antigravity, and Gemini CLI

Providers: google-vertex , google-antigravity , google-gemini-cli

Auth: Vertex uses gcloud ADC; Antigravity/Gemini CLI use their respective auth flows

Antigravity OAuth is shipped as a bundled plugin ( google-antigravity-auth , disabled by default).

Enable: openclaw plugins enable google-antigravity-auth

Login: openclaw models auth login --provider google-antigravity --set-default

Gemini CLI OAuth is shipped as a bundled plugin ( google-gemini-cli-auth , disabled by default).

Enable: openclaw plugins enable google-gemini-cli-auth

Login: openclaw models auth login --provider google-gemini-cli --set-default

Note: you do **not** paste a client id or secret into openclaw.json . The CLI login flow stores tokens in auth profiles on the gateway host.

## Z.AI (GLM)



Provider: zai

Auth: ZAI\_API\_KEY

Example model: zai/glm-4.7

CLI: openclaw onboard --auth-choice zai-api-key

Aliases: z.ai/\* and z-ai/\* normalize to zai/\*

## Vercel AI Gateway

Provider: vercel-ai-gateway

Auth: AI\_GATEWAY\_API\_KEY

Example model: vercel-ai-gateway/anthropic/clause-opus-4.6

CLI: openclaw onboard --auth-choice ai-gateway-api-key

## Other built-in providers

OpenRouter: openrouter ( OPENROUTER\_API\_KEY )

Example model: openrouter/anthropic/clause-sonnet-4-5

xAI: xai ( XAI\_API\_KEY )

Groq: groq ( GROQ\_API\_KEY )

Cerebras: cerebras ( CEREBRAS\_API\_KEY )

GLM models on Cerebras use ids zai-glm-4.7 and zai-glm-4.6 .

OpenAI-compatible base URL: <https://api.cerebras.ai/v1> .

Mistral: mistral ( MISTRAL\_API\_KEY )

GitHub Copilot: github-copilot ( COPILOT\_GITHUB\_TOKEN / GH\_TOKEN / GITHUB\_TOKEN )

Hugging Face Inference: huggingface ( HUGGINGFACE\_HUB\_TOKEN or HF\_TOKEN ) – OpenAI-compatible router; example model: huggingface/deepseek-ai/DeepSeek-R1 ; CLI: openclaw onboard --auth-choice huggingface-api-key . See [Hugging Face \(Inference\)](#).

## Providers via `models.providers` (custom/base URL)



Use `models.providers` (or `models.json`) to add **custom** providers or OpenAI/Anthropic-compatible proxies.

### Moonshot AI (Kimi)

Moonshot uses OpenAI-compatible endpoints, so configure it as a custom provider:

Provider: `moonshot`

Auth: `MOONSHOT_API_KEY`

Example model: `moonshot/kimi-k2.5`

Kimi K2 model IDs:

`moonshot/kimi-k2.5`

`moonshot/kimi-k2-0905-preview`

`moonshot/kimi-k2-turbo-preview`

`moonshot/kimi-k2-thinking`

`moonshot/kimi-k2-thinking-turbo`



```

agents: {
  defaults: { model: { primary: "moonshot/kimi-k2.5" } },
},
models: {
  mode: "merge",
  providers: {
    moonshot: {
      baseUrl: "https://api.moonshot.ai/v1",
      apiKey: "${MOONSHOT_API_KEY}",
      api: "openai-completions",
      models: [{ id: "kimi-k2.5", name: "Kimi K2.5" }],
    },
  },
},
}

```

## Kimi Coding

Kimi Coding uses Moonshot AI's Anthropic-compatible endpoint:

Provider: kimi-coding

Auth: KIMI\_API\_KEY

Example model: kimi-coding/k2p5

```
{
  env: { KIMI_API_KEY: "sk-..." },
  agents: {
    defaults: { model: { primary: "kimi-coding/k2p5" } },
  },
}
```

## Qwen OAuth (free tier)

Qwen provides OAuth access to Qwen Coder + Vision via a device-code flow. Enable the bundled plugin, then log in:

```
openclaw plugins enable qwen-portal-auth  
openclaw models auth login --provider qwen-portal --set-default
```

Model refs:

qwen-portal/coder-model

qwen-portal/vision-model

See [for setup details and notes.](#)

## Synthetic

Synthetic provides Anthropic-compatible models behind the `synthetic` provider:

Provider: `synthetic`

Auth: `SYNTHETIC_API_KEY`

Example model: `synthetic/hf:MiniMaxAI/MiniMax-M2.1`

CLI: `openclaw onboard --auth-choice synthetic-api-key`



```
agents: {
    defaults: { model: { primary: "synthetic/hf:MiniMaxAI/MiniMax-M2.1" } },
},
models: {
    mode: "merge",
    providers: {
        synthetic: {
            baseUrl: "https://api.synthetic.new/anthropic",
            apiKey: "${SYNTHETIC_API_KEY}",
            api: "anthropic-messages",
            models: [{ id: "hf:MiniMaxAI/MiniMax-M2.1", name: "MiniMax M2.1" }],
        },
    },
},
}
```

## MinMax

MinMax is configured via `models.providers` because it uses custom endpoints:

```
MinMax (Anthropic-compatible): --auth-choice minimax-api  
Auth: MINIMAX_API_KEY
```

See [here](#) for setup details, model options, and config snippets.

## Ollama

Ollama is a local LLM runtime that provides an OpenAI-compatible API:

```
Provider: ollama  
Auth: None required (local server)  
Example model: ollama/llama3.3
```

Installation: <https://ollama.ai>



```
# Install Ollama, then pull a model:  
ollama pull llama3.3
```

```
{  
  agents: {  
    defaults: { model: { primary: "ollama/llama3.3" } },  
  },  
}
```

Ollama is automatically detected when running locally at <http://127.0.0.1:11434/v1>. See [the documentation](#) for model recommendations and custom configuration.

## vLLM

vLLM is a local (or self-hosted) OpenAI-compatible server:

Provider: vllm

Auth: Optional (depends on your server)

Default base URL: <http://127.0.0.1:8000/v1>

To opt in to auto-discovery locally (any value works if your server doesn't enforce auth):

```
export VLLM_API_KEY="vllm-local"
```

Then set a model (replace with one of the IDs returned by `/v1/models`):



```
agents: {  
    defaults: { model: { primary: "vllm/your-model-id" } },  
},  
}
```

See [for details.](#)

## Local proxies (LM Studio, vLLM, LiteLLM, etc.)

Example (OpenAI-compatible):



```

agents: {
  defaults: {
    model: { primary: "lmstudio/minimax-m2.1-gs32" },
    models: { "lmstudio/minimax-m2.1-gs32": { alias: "Minimax" } },
  },
  models: {
    providers: {
      lmstudio: {
        baseUrl: "http://localhost:1234/v1",
        apiKey: "LMSTUDIO_KEY",
        api: "openai-completions",
        models: [
          {
            id: "minimax-m2.1-gs32",
            name: "MiniMax M2.1",
            reasoning: false,
            input: ["text"],
            cost: { input: 0, output: 0, cacheRead: 0, cacheWrite: 0 },
            contextWindow: 200000,
            maxTokens: 8192,
          },
        ],
      },
    },
  },
}

```

## Notes:

For custom providers, reasoning , input , cost , contextWindow , and maxTokens are optional. When omitted, OpenClaw defaults to:

```

reasoning: false

input: ["text"]

cost: { input: 0, output: 0, cacheRead: 0, cacheWrite: 0 }

```



contextWindow: 200000

maxTokens: 8192

Recommended:, set explicit values that match your proxy/model limits.

## CLI examples

```
openclaw onboard --auth-choice opencode-zen  
openclaw models set opencode/clause-opus-4-6  
openclaw models list
```

See also:

for full configuration examples.

[Models CLI](#)

[Model Failover](#)

Powered by [mintlify](#)