



☰ Providers > Moonshot AI

## Providers

# Moonshot AI

Moonshot provides the Kimi API with OpenAI-compatible endpoints. Configure the provider and set the default model to `moonshot/kimi-k2.5`, or use Kimi Coding with `kimi-coding/k2p5`.

Current Kimi K2 model IDs:

```
kimi-k2.5  
kimi-k2-0905-preview  
kimi-k2-turbo-preview  
kimi-k2-thinking  
kimi-k2-thinking-turbo
```

```
openclaw onboard --auth-choice moonshot-api-key
```

Kimi Coding:

```
openclaw onboard --auth-choice kimi-code-api-key
```

Note: Moonshot and Kimi Coding are separate providers. Keys are not interchangeable, endpoints differ, and model refs differ (Moonshot uses `moonshot/...`, Kimi Coding uses `kimi-coding/...`).

# Config snippet (Moonshot API)

&gt;



```
env: { MOONSHOT_API_KEY: "sk-..." },
agents: {
  defaults: {
    model: { primary: "moonshot/kimi-k2.5" },
    models: {
      // moonshot-kimi-k2-aliases:start
      "moonshot/kimi-k2.5": { alias: "Kimi K2.5" },
      "moonshot/kimi-k2-0905-preview": { alias: "Kimi K2" },
      "moonshot/kimi-k2-turbo-preview": { alias: "Kimi K2 Turbo" },
      "moonshot/kimi-k2-thinking": { alias: "Kimi K2 Thinking" },
      "moonshot/kimi-k2-thinking-turbo": { alias: "Kimi K2 Thinking Turbo" },
      // moonshot-kimi-k2-aliases:end
    },
  },
},
models: {
  mode: "merge",
  providers: {
    moonshot: {
      baseUrl: "https://api.moonshot.ai/v1",
      apiKey: "${MOONSHOT_API_KEY}",
      api: "openai-completions",
      models: [
        // moonshot-kimi-k2-models:start
        {
          id: "kimi-k2.5",
          name: "Kimi K2.5",
          reasoning: false,
          input: ["text"],
          cost: { input: 0, output: 0, cacheRead: 0, cacheWrite: 0 },
          contextWindow: 256000,
          maxTokens: 8192,
        },
        {
          id: "kimi-k2-0905-preview",
          name: "Kimi K2 0905 Preview",
          reasoning: false,
          input: ["text"],
          cost: { input: 0, output: 0, cacheRead: 0, cacheWrite: 0 },
        }
      ]
    }
  }
}
```



```
        contextWindow: 256000,
        maxTokens: 8192,
    },
    {
        >
        id: "kimi-k2-turbo-preview",
        name: "Kimi K2 Turbo",
        reasoning: false,
        input: ["text"],
        cost: { input: 0, output: 0, cacheRead: 0, cacheWrite: 0 },
        contextWindow: 256000,
        maxTokens: 8192,
    },
    {
        >
        id: "kimi-k2-thinking",
        name: "Kimi K2 Thinking",
        reasoning: true,
        input: ["text"],
        cost: { input: 0, output: 0, cacheRead: 0, cacheWrite: 0 },
        contextWindow: 256000,
        maxTokens: 8192,
    },
    {
        >
        id: "kimi-k2-thinking-turbo",
        name: "Kimi K2 Thinking Turbo",
        reasoning: true,
        input: ["text"],
        cost: { input: 0, output: 0, cacheRead: 0, cacheWrite: 0 },
        contextWindow: 256000,
        maxTokens: 8192,
    },
    // moonshot-kimi-k2-models:end
],
},
},
},
}
```

## Kimi Coding



```
env: { KIMI_API_KEY: "sk-..." },
agents: {
  defaults: {
    model: { primary: "kimi-coding/k2p5" },
    models: {
      "kimi-coding/k2p5": { alias: "Kimi K2.5" },
    },
  },
},
}
```

## Notes

Moonshot model refs use `moonshot/<modelId>`. Kimi Coding model refs use `kimi-coding/<modelId>`.

Override pricing and context metadata in `models.providers` if needed.

If Moonshot publishes different context limits for a model, adjust `contextWindow` accordingly.

Use `https://api.moonshot.ai/v1` for the international endpoint, and `https://api.moonshot.cn/v1` for the China endpoint.

◀ Vercel AI Gateway

Minimax ▶

Powered by [mintlify](#)