



Providers

Synthetic

Synthetic exposes Anthropic-compatible endpoints. OpenClaw registers it as the `synthetic` provider and uses the Anthropic Messages API.

Quick setup

1. Set `SYNTHETIC_API_KEY` (or run the wizard below).
2. Run onboarding:

```
openclaw onboard --auth-choice synthetic-api-key
```

The default model is set to:

```
synthetic/hf:MiniMaxAI/MiniMax-M2.1
```

Config example



```
env: { SYNTHETIC_API_KEY: "sk-..." },
agents: {
  defaults: {
    model: { primary: "synthetic/hf:MiniMaxAI/MiniMax-M2.1" },
    models: { "synthetic/hf:MiniMaxAI/MiniMax-M2.1": { alias: "MiniMax M2.1" } }
  },
},
models: {
  mode: "merge",
  providers: {
    synthetic: {
      baseUrl: "https://api.synthetic.new/anthropic",
      apiKey: "${SYNTHETIC_API_KEY}",
      api: "anthropic-messages",
      models: [
        {
          id: "hf:MiniMaxAI/MiniMax-M2.1",
          name: "MiniMax M2.1",
          reasoning: false,
          input: ["text"],
          cost: { input: 0, output: 0, cacheRead: 0, cacheWrite: 0 },
          contextWindow: 192000,
          maxTokens: 65536,
        },
      ],
    },
  },
},
}
```

Note: OpenClaw's Anthropic client appends `/v1` to the base URL, so use `https://api.synthetic.new/anthropic` (not `/anthropic/v1`). If Synthetic changes its base URL, override `models.providers.synthetic.baseUrl`.

Model catalog

All models below use cost 0 (input/output/cache).



Model ID	Context window	Max tokens	Reasoning	Input
hf:MiniMaxAI/MiniM ax-M2.1	192000	65536	false	text
hf:moonshotai/Kimi -K2-Thinking	256000	8192	true	text
hf:zai-org/GLM- 4.7	198000	128000	false	text
hf:deepseek- ai/DeepSeek-R1- 0528	128000	8192	false	text
hf:deepseek- ai/DeepSeek-V3- 0324	128000	8192	false	text
hf:deepseek- ai/DeepSeek-V3.1	128000	8192	false	text
hf:deepseek- ai/DeepSeek-V3.1- Terminus	128000	8192	false	text
hf:deepseek- ai/DeepSeek-V3.2	159000	8192	false	text
hf:meta- llama/Llama-3.3- 70B-Instruct	128000	8192	false	text
hf:meta- llama/Llama-4- Maverick-17B-128E- Instruct-FP8	524000	8192	false	text
hf:moonshotai/Kimi -K2-Instruct-0905	256000	8192	false	text
hf:openai/gpt-oss- 120b	128000	8192	false	text
hf:Qwen/Qwen3- 235B-A22B- Instruct-2507	256000	8192	false	text

Model ID	Context window	Max tokens	Reasoning	Input
hf:Qwen/Qwen3-Coder-480B-A35B-Instruct	256000	8192	false	text
hf:Qwen/Qwen3-VL-235B-A22B-Instruct	250000	8192	false	text + image
hf:zai-org/GLM-4.5	128000	128000	false	text
hf:zai-org/GLM-4.6	198000	128000	false	text
hf:deepseek-ai/DeepSeek-V3	128000	8192	false	text
hf:Qwen/Qwen3-235B-A22B-Thinking-2507	256000	8192	true	text

Notes

Model refs use `synthetic/<modelId>`.

If you enable a model allowlist (`agents.defaults.models`), add every model you plan to use.

See [Model providers](#) for provider rules.