



macOS companion app > **Voice Overlay**

macOS companion app

Voice Overlay

Audience: macOS app contributors. Goal: keep the voice overlay predictable when wake-word and push-to-talk overlap.

Current intent

If the overlay is already visible from wake-word and the user presses the hotkey, the hotkey session *adopts* the existing text instead of resetting it. The overlay stays up while the hotkey is held. When the user releases: send if there is trimmed text, otherwise dismiss.

Wake-word alone still auto-sends on silence; push-to-talk sends immediately on release.

Implemented (Dec 9, 2025)

Overlay sessions now carry a token per capture (wake-word or push-to-talk). Partial/final/send/dismiss/level updates are dropped when the token doesn't match, avoiding stale callbacks.

Push-to-talk adopts any visible overlay text as a prefix (so pressing the hotkey while the wake overlay is up keeps the text and appends new speech). It waits up to 1.5s for a final transcript before falling back to the current text.

Chime/overlay logging is emitted at `info` in categories `voicewake.overlay`, `voicewake.ptt`, and `voicewake.chime` (session start,



partial, final, send, dismiss, chime reason).

Next steps

>

1. VoiceSessionCoordinator (actor)

Owns exactly one `VoiceSession` at a time.

API (token-based): `beginWakeCapture`, `beginPushToTalk`,
`updatePartial`, `endCapture`, `cancel`, `applyCooldown`.

Drops callbacks that carry stale tokens (prevents old
recognizers from reopening the overlay).

2. VoiceSession (model)

Fields: `token`, `source` (`wakeWord|pushToTalk`),
committed/volatile text, chime flags, timers (auto-send, idle),
`overlayMode` (`display|editing|sending`), cooldown deadline.

3. Overlay binding

`VoiceSessionPublisher` (`ObservableObject`) mirrors the active
session into SwiftUI.

`VoiceWakeOverlayView` renders only via the publisher; it never
mutates global singletons directly.

Overlay user actions (`sendNow`, `dismiss`, `edit`) call back into
the coordinator with the session token.

4. Unified send path

On `endCapture`: if trimmed text is empty → dismiss; else
`performSend(session:)` (plays send chime once, forwards,
dismisses).

Push-to-talk: no delay; wake-word: optional delay for auto-
send.

Apply a short cooldown to the wake runtime after push-to-talk
finishes so wake-word doesn't immediately retrigger.

5. Logging



Coordinator emits `.info` logs in subsystem `bot.molt`, categories `voicewake.overlay` and `voicewake.chime`.

Key events: `session_started`, `adopted_by_push_to_talk`, `partial`, `finalized`, `send`, `dismiss`, `cancel`, `cooldown`.

Debugging checklist

Stream logs while reproducing a sticky overlay:

```
sudo log stream --predicate 'subsystem == "bot.molt" AND category
```

Verify only one active session token; stale callbacks should be dropped by the coordinator.

Ensure push-to-talk release always calls `endCapture` with the active token; if text is empty, expect `dismiss` without chime or send.

Migration steps (suggested)

1. Add `VoiceSessionCoordinator`, `VoiceSession`, and `VoiceSessionPublisher`.
2. Refactor `VoiceWakeRuntime` to create/update/end sessions instead of touching `VoiceWakeOverlayController` directly.
3. Refactor `VoicePushToTalk` to adopt existing sessions and call `endCapture` on release; apply runtime cooldown.
4. Wire `VoiceWakeOverlayController` to the publisher; remove direct calls from runtime/PTT.
5. Add integration tests for session adoption, cooldown, and empty-text dismissal.



Powered by **mintlify**

>
