



☰ Media and devices > **Audio and Voice Notes**

Media and devices

Audio and Voice Notes

What works

Media understanding (audio): If audio understanding is enabled (or auto-detected), OpenClaw:

1. Locates the first audio attachment (local path or URL) and downloads it if needed.
2. Enforces `maxBytes` before sending to each model entry.
3. Runs the first eligible model entry in order (provider or CLI).
4. If it fails or skips (size/timeout), it tries the next entry.
5. On success, it replaces `Body` with an `[Audio]` block and sets `{{Transcript}}`.

Command parsing: When transcription succeeds, `CommandBody` / `RawBody` are set to the transcript so slash commands still work.

Verbose logging: In `--verbose`, we log when transcription runs and when it replaces the body.

Auto-detection (default)

If you **don't configure models** and `tools.media.audio.enabled` is **not** set to `false`, OpenClaw auto-detects in this order and stops at the first working option:

1. **Local CLIs** (if installed)



`sherpa-onnx-offline` (requires `SHERPA_ONNX_MODEL_DIR` with encoder/decoder/joiner/tokens)

`whisper-cli` (from `whisper-cpp` ; uses `WHISPER_CPP_MODEL` or the bundled tiny model)

`whisper` (Python CLI; downloads models automatically)

2. **Gemini CLI** (`gemini`) using `read_many_files`

3. **Provider keys** (OpenAI → Groq → Deepgram → Google)

To disable auto-detection, set `tools.media.audio.enabled: false`. To customize, set `tools.media.audio.models`. Note: Binary detection is best-effort across macOS/Linux/Windows; ensure the CLI is on `PATH` (we expand `~`), or set an explicit CLI model with a full command path.

Config examples

Provider + CLI fallback (OpenAI + Whisper CLI)



```
tools: {
  media: {
    audio: {
      enabled: true,
      maxBytes: 20971520,
      models: [
        { provider: "openai", model: "gpt-4o-mini-transcribe" },
        {
          type: "cli",
          command: "whisper",
          args: ["--model", "base", "{{MediaPath}}"],
          timeoutSeconds: 45,
        },
      ],
    },
  },
}
```

Provider-only with scope gating

```
{
  tools: {
    media: {
      audio: {
        enabled: true,
        scope: {
          default: "allow",
          rules: [{ action: "deny", match: { chatType: "group" } }],
        },
        models: [{ provider: "openai", model: "gpt-4o-mini-transcribe" }],
      },
    },
  },
}
```

Provider-only (Deepgram)

```
{
  tools: {
    media: {
      audio: {
        enabled: true,
        models: [{ provider: "deepgram", model: "nova-3" }],
      },
    },
  },
}
```

Notes & limits

Provider auth follows the standard model auth order (auth profiles, env vars, `models.providers.*.apiKey`).

Deepgram picks up `DEEPGRAM_API_KEY` when `provider: "deepgram"` is used.

Deepgram setup details:

Audio providers can override `baseUrl`, `headers`, and `providerOptions` via `tools.media.audio`.

Default size cap is 20MB (`tools.media.audio.maxBytes`). Oversize audio is skipped for that model and the next entry is tried.

Default `maxChars` for audio is `unset` (full transcript). Set `tools.media.audio.maxChars` or per-entry `maxChars` to trim output.

OpenAI auto default is `gpt-4o-mini-transcribe`; set `model: "gpt-4o-transcribe"` for higher accuracy.

Use `tools.media.audio.attachments` to process multiple voice notes (`mode: "all" + maxAttachments`).

Transcript is available to templates as `{Transcript}`.

CLI stdout is capped (5MB); keep CLI output concise.

Mention Detection in Groups



When `requireMention: true` is set for a group chat, OpenClaw now transcribes audio **before** checking for mentions. This allows voice notes to be processed even when they contain mentions.

How it works:

1. If a voice message has no text body and the group requires mentions, OpenClaw performs a “preflight” transcription.
2. The transcript is checked for mention patterns (e.g., `@BotName`, emoji triggers).
3. If a mention is found, the message proceeds through the full reply pipeline.
4. The transcript is used for mention detection so voice notes can pass the mention gate.

Fallback behavior:

If transcription fails during preflight (timeout, API error, etc.), the message is processed based on text-only mention detection.

This ensures that mixed messages (text + audio) are never incorrectly dropped.

Example: A user sends a voice note saying “Hey @Claude, what’s the weather?” in a Telegram group with `requireMention: true`. The voice note is transcribed, the mention is detected, and the agent replies.

Gotchas

Scope rules use first-match wins. `chatType` is normalized to `direct`, `group`, or `room`.

Ensure your CLI exits `0` and prints plain text; JSON needs to be massaged via `jq -r .text`.

 Keep timeouts reasonable (`timeoutSeconds` , default 60s) to avoid blocking the reply queue.

Preflight transcription only processes the **first** audio attachment for mention detection. Additional audio is processed during the main media understanding phase.

[◀ Image and Media Support](#)

[Camera Capture ▶](#)

Powered by [mintlify](#)