



☰ Built-in tools > Exec Tool

Built-in tools

Exec Tool

Run shell commands in the workspace. Supports foreground + background execution via `process`. If `process` is disallowed, `exec` runs synchronously and ignores `yieldMs` / `background`. Background sessions are scoped per agent; `process` only sees sessions from the same agent.

) Parameters

```
command (required)

workdir (defaults to cwd)

env (key/value overrides)

yieldMs (default 10000): auto-background after delay

background (bool): background immediately

timeout (seconds, default 1800): kill on expiry

pty (bool): run in a pseudo-terminal when available (TTY-only
CLIs, coding agents, terminal UIs)

host ( sandbox | gateway | node ): where to execute

security ( deny | allowlist | full ): enforcement mode for
gateway / node

ask ( off | on-miss | always ): approval prompts for gateway / node

node (string): node id/name for host=node

elevated (bool): request elevated mode (gateway host);
security=full is only forced when elevated resolves to full
```

Notes:



host defaults to sandbox .

elevated is ignored when sandboxing is off (exec already runs on the host).

gateway / node approvals are controlled by `~/.openclaw/exec-approvals.json` .

node requires a paired node (companion app or headless node host).

If multiple nodes are available, set `exec.node` or `tools.exec.node` to select one.

On non-Windows hosts, exec uses SHELL when set; if SHELL is fish , it prefers bash (or sh) from PATH to avoid fish-incompatible scripts, then falls back to SHELL if neither exists.

Host execution (gateway / node) rejects env.PATH and loader overrides (LD_* / DYLD_*) to prevent binary hijacking or injected code.

Important: sandboxing is off by default. If sandboxing is off, host=sandbox runs directly on the gateway host (no container) and does not require approvals. To require approvals, run with host=gateway and configure exec approvals (or enable sandboxing).

Config

`tools.exec.notifyOnExit` (default: true): when true, backgrounded exec sessions enqueue a system event and request a heartbeat on exit.

`tools.exec.approvalRunningNoticeMs` (default: 10000): emit a single “running” notice when an approval-gated exec runs longer than this (0 disables).

`tools.exec.host` (default: sandbox)

`tools.exec.security` (default: deny for sandbox, allowlist for gateway + node when unset)

 tools.exec.ask (default: on-miss)
 tools.exec.node (default: unset)
 tools.exec.pathPrepend : list of directories to prepend to PATH for exec runs (gateway + sandbox only).
 tools.exec.safeBins : stdin-only safe binaries that can run without explicit allowlist entries. For behavior details, see [Safe bins](#).

Example:

```
{
  tools: {
    exec: {
      pathPrepend: ["~/bin", "/opt/oss/bin"],
    },
  },
}
```

PATH handling

host=gateway : merges your login-shell PATH into the exec environment. env.PATH overrides are rejected for host execution. The daemon itself still runs with a minimal PATH :

macOS: /opt/homebrew/bin , /usr/local/bin , /usr/bin , /bin

Linux: /usr/local/bin , /usr/bin , /bin

host=sandbox : runs sh -lc (login shell) inside the container, so /etc/profile may reset PATH . OpenClaw prepends env.PATH after profile sourcing via an internal env var (no shell interpolation); tools.exec.pathPrepend applies here too.

host=node : only non-blocked env overrides you pass are sent to the node. env.PATH overrides are rejected for host execution and ignored by node hosts. If you need additional PATH entries on a node, configure the node host service environment (systemd/launchd) or install tools in standard locations.

Per-agent node binding (use the agent list index in config):



```
openclaw config get agents.list
openclaw config set agents.list[0].tools.exec.node "node-id-or-name"
```

Control UI: the Nodes tab includes a small “Exec node binding” panel for the same settings.

Session overrides (/exec)

Use `/exec` to set **per-session** defaults for `host`, `security`, `ask`, and `node`. Send `/exec` with no arguments to show the current values.

Example:

```
/exec host=gateway security=allowlist ask=on-miss node=mac-1
```

Authorization model

`/exec` is only honored for **authorized senders** (channel allowlists/pairing plus `commands.useAccessGroups`). It updates **session state only** and does not write config. To hard-disable exec, deny it via tool policy (`tools.deny: ["exec"]` or per-agent). Host approvals still apply unless you explicitly set `security=full` and `ask=off`.

Exec approvals (companion app / node host)

Sandboxed agents can require per-request approval before `exec` runs on the gateway or node host. See [the companion app](#) for the policy, allowlist, and UI flow.

When approvals are required, the exec tool returns immediately with `status: "approval-pending"` and an approval id. Once approved (or denied /

timed out), the Gateway emits system events (Exec finished / Exec denied). If the command is still running after tools.exec.approvalRunningNoticeMs , a single Exec running notice is emitted.

Allowlist + safe bins

Allowlist enforcement matches **resolved binary paths only** (no basename matches). When security=allowlist , shell commands are auto-allowed only if every pipeline segment is allowlisted or a safe bin. Chaining (; , && , ||) and redirections are rejected in allowlist mode unless every top-level segment satisfies the allowlist (including safe bins). Redirections remain unsupported.

Examples

Foreground:

```
{ "tool": "exec", "command": "ls -la" }
```

Background + poll:

```
{"tool":"exec","command":"npm run build","yieldMs":1000}
{"tool":"process","action":"poll","sessionId":<id>"}
```

Send keys (tmux-style):

```
{"tool":"process","action":"send-keys","sessionId":<id>,"keys":["End"]
 {"tool":"process","action":"send-keys","sessionId":<id>,"keys":["C-c"]}
 {"tool":"process","action":"send-keys","sessionId":<id>,"keys":["Up","Up","Enter"]}
```

Submit (send CR only):



```
{ "tool": "process", "action": "submit", "sessionId": "<id>" }
```

Paste (bracketed by default):

```
{ "tool": "process", "action": "paste", "sessionId": "<id>", "text": '
```

apply_patch (experimental)

`apply_patch` is a subtool of `exec` for structured multi-file edits.

Enable it explicitly:

```
{
  tools: {
    exec: {
      applyPatch: { enabled: true, workspaceOnly: true, allowModels: ["gpt-5.2"] },
    },
  },
}
```

Notes:

Only available for OpenAI/OpenAI Codex models.

Tool policy still applies; `allow: ["exec"]` implicitly allows `apply_patch`.

Config lives under `tools.exec.applyPatch`.

`tools.exec.applyPatch.workspaceOnly` defaults to `true` (workspace-contained). Set it to `false` only if you intentionally want `apply_patch` to write/delete outside the workspace directory.

< LLM Task



Web Tools >

Powered by mintlify

>