



≡ Built-in tools > Lobster

Built-in tools

Lobster

Typed workflow runtime for OpenClaw – composable pipelines with approval gates.

Lobster is a workflow shell that lets OpenClaw run multi-step tool sequences as a single, deterministic operation with explicit approval checkpoints.

Hook

Your assistant can build the tools that manage itself. Ask for a workflow, and 30 minutes later you have a CLI plus pipelines that run as one call. Lobster is the missing piece: deterministic pipelines, explicit approvals, and resumable state.

Why

Today, complex workflows require many back-and-forth tool calls. Each call costs tokens, and the LLM has to orchestrate every step. Lobster moves that orchestration into a typed runtime:

One call instead of many: OpenClaw runs one Lobster tool call and gets a structured result.

Approvals built in: Side effects (send email, post comment) halt the workflow until explicitly approved.



Resumable: Halted workflows return a token; approve and resume without re-running everything.

>

Why a DSL instead of plain programs?

Lobster is intentionally small. The goal is not “a new language,” it’s a predictable, AI-friendly pipeline spec with first-class approvals and resume tokens.

Approve/resume is built in: A normal program can prompt a human, but it can’t pause and resume with a durable token without you inventing that runtime yourself.

Determinism + auditability: Pipelines are data, so they’re easy to log, diff, replay, and review.

Constrained surface for AI: A tiny grammar + JSON piping reduces “creative” code paths and makes validation realistic.

Safety policy baked in: Timeouts, output caps, sandbox checks, and allowlists are enforced by the runtime, not each script.

Still programmable: Each step can call any CLI or script. If you want JS/TS, generate `.lobster` files from code.

How it works

OpenClaw launches the local `lobster` CLI in **tool mode** and parses a JSON envelope from stdout. If the pipeline pauses for approval, the tool returns a `resumeToken` so you can continue later.

Pattern: small CLI + JSON pipes + approvals

Build tiny commands that speak JSON, then chain them into a single Lobster call. (Example command names below – swap in your own.)

```
 inbox list --json
inbox categorize --json
inbox apply --json
    >
```

```
{
  "action": "run",
  "pipeline": "exec --json --shell 'inbox list --json' | exec --stdin json --shell 'categorize --json' | exec --stdin json --shell 'apply --json' > /tmp/inbox.json"
  "timeoutMs": 30000
}
```

If the pipeline requests approval, resume with the token:

```
{
  "action": "resume",
  "token": "<resumeToken>",
  "approve": true
}
```

AI triggers the workflow; Lobster executes the steps. Approval gates keep side effects explicit and auditable.

Example: map input items into tool calls:

```
gog.gmail.search --query 'newer_than:1d' \
| openclaw.invoke --tool message --action send --each --item-key message --args
```

JSON-only LLM steps (llm-task)

For workflows that need a **structured LLM step**, enable the optional `llm-task` plugin tool and call it from Lobster. This keeps the workflow deterministic while still letting you classify/summarize/draft with a model.

Enable the tool:



```
{  
  "plugins": {  
    "entries": [  
      "llm-task": { "enabled": true }  
    ]  
  },  
  "agents": {  
    "list": [  
      {  
        "id": "main",  
        "tools": { "allow": ["llm-task"] }  
      }  
    ]  
  }  
}
```

Use it in a pipeline:

```
openclaw.invoke --tool llm-task --action json --args-json '{  
  "prompt": "Given the input email, return intent and draft.",  
  "input": { "subject": "Hello", "body": "Can you help?" },  
  "schema": {  
    "type": "object",  
    "properties": {  
      "intent": { "type": "string" },  
      "draft": { "type": "string" }  
    },  
    "required": [ "intent", "draft" ],  
    "additionalProperties": false  
  }  
}'
```

See [the configuration documentation](#) for details and configuration options.

Workflow files (.lobster)



Lobster can run YAML/JSON workflow files with `name`, `args`, `steps`, `env`, `condition`, and `approval` fields. In OpenClaw tool calls, set `pipeline` to the file path.

```
name: inbox-triage
args:
  tag:
    default: "family"
steps:
  - id: collect
    command: inbox list --json
  - id: categorize
    command: inbox categorize --json
    stdin: $collect.stdout
  - id: approve
    command: inbox apply --approve
    stdin: $categorize.stdout
    approval: required
  - id: execute
    command: inbox apply --execute
    stdin: $categorize.stdout
    condition: $approve.approved
```

Notes:

`stdin: $step.stdout` and `stdin: $step.json` pass a prior step's output.
`condition` (or `when`) can gate steps on `$step.approved`.

Install Lobster

Install the Lobster CLI on the `same host` that runs the OpenClaw Gateway (see the), and ensure `lobster` is on `PATH`. If you want to use a custom binary location, pass an `absolute lobsterPath` in the tool call.

Enable the tool



Lobster is an **optional** plugin tool (not enabled by default).

>
Recommended (additive, safe):

```
{  
  "tools": {  
    "alsoAllow": ["lobster"]  
  }  
}
```

Or per-agent:

```
{  
  "agents": {  
    "list": [  
      {  
        "id": "main",  
        "tools": {  
          "alsoAllow": ["lobster"]  
        }  
      }  
    ]  
  }  
}
```

Avoid using `tools.allow: ["lobster"]` unless you intend to run in restrictive allowlist mode.

Note: allowlists are opt-in for optional plugins. If your allowlist only names plugin tools (like `lobster`), OpenClaw keeps core tools enabled. To restrict core tools, include the core tools or groups you want in the allowlist too.

Example: Email triage



Without Lobster:

>

```
User: "Check my email and draft replies"
→ openclaw calls gmail.list
→ LLM summarizes
→ User: "draft replies to #2 and #5"
→ LLM drafts
→ User: "send #2"
→ openclaw calls gmail.send
(repeat daily, no memory of what was triaged)
```

With Lobster:

```
{
  "action": "run",
  "pipeline": "email.triage --limit 20",
  "timeoutMs": 30000
}
```

Returns a JSON envelope (truncated):

```
{
  "ok": true,
  "status": "needs_approval",
  "output": [{ "summary": "5 need replies, 2 need action" }],
  "requiresApproval": {
    "type": "approval_request",
    "prompt": "Send 2 draft replies?",
    "items": [],
    "resumeToken": "..."
  }
}
```

User approves → resume:



```
{  
  "action": "resume",  
  "token": "<resumeToken>",  
  "approve": true  
}
```

One workflow. Deterministic. Safe.

Tool parameters

run

Run a pipeline in tool mode.

```
{  
  "action": "run",  
  "pipeline": "gog.gmail.search --query 'newer_than:1d' | email.triage",  
  "cwd": "/path/to/workspace",  
  "timeoutMs": 30000,  
  "maxStdoutBytes": 512000  
}
```

Run a workflow file with args:

```
{  
  "action": "run",  
  "pipeline": "/path/to/inbox-triage.lobster",  
  "argsJson": "{\"tag\": \"family\"}"  
}
```

resume

Continue a halted workflow after approval.



```
{  
  "action": "resume",  
  "token": "<resumeToken>",  
  "approve": true  
}
```

Optional inputs

`lobsterPath` : Absolute path to the Lobster binary (omit to use `PATH`).

`cwd` : Working directory for the pipeline (defaults to the current process working directory).

`timeoutMs` : Kill the subprocess if it exceeds this duration (default: 20000).

`maxStdoutBytes` : Kill the subprocess if stdout exceeds this size (default: 512000).

`argsJson` : JSON string passed to `lobster run --args-json` (workflow files only).

Output envelope

Lobster returns a JSON envelope with one of three statuses:

`ok` → finished successfully

`needs_approval` → paused; `requiresApproval.resumeToken` is required to resume

`cancelled` → explicitly denied or cancelled

The tool surfaces the envelope in both `content` (pretty JSON) and `details` (raw object).

Approvals



If `requiresApproval` is present, inspect the prompt and decide:

>
`approve: true` → resume and continue side effects

`approve: false` → cancel and finalize the workflow

Use `approve --preview-from-stdin --limit N` to attach a JSON preview to approval requests without custom jq/heredoc glue. Resume tokens are now compact: Lobster stores workflow resume state under its state dir and hands back a small token key.

OpenProse

OpenProse pairs well with Lobster: use `/prose` to orchestrate multi-agent prep, then run a Lobster pipeline for deterministic approvals. If a Prose program needs Lobster, allow the `lobster` tool for sub-agents via `tools.subagents.tools`. See [OpenProse](#).

Safety

Local subprocess only – no network calls from the plugin itself.

No secrets – Lobster doesn't manage OAuth; it calls OpenClaw tools that do.

Sandbox-aware – disabled when the tool context is sandboxed.

Hardened – `lobsterPath` must be absolute if specified; timeouts and output caps enforced.

Troubleshooting

`lobster subprocess timed out` → increase `timeoutMs`, or split a long pipeline.



lobster output exceeded maxStdoutBytes → raise maxStdoutBytes or reduce output size.

lobster returned invalid JSON → ensure the pipeline runs in tool mode
and prints only JSON.

lobster failed (code ...) → run the same pipeline in a terminal to inspect stderr.

Learn more

[Plugins](#)

[Plugin tool authoring](#)

Case study: community workflows

One public example: a “second brain” CLI + Lobster pipelines that manage three Markdown vaults (personal, partner, shared). The CLI emits JSON for stats, inbox listings, and stale scans; Lobster chains those commands into workflows like `weekly-review`, `inbox-triage`, `memory-consolidation`, and `shared-task-sync`, each with approval gates. AI handles judgment (categorization) when available and falls back to deterministic rules when not.

Thread: <https://x.com/plattenschieber/status/2014508656335770033>

Repo: <https://github.com/bloomedai/brain-cli>

◀ Tools

LLM Task ▶

Powered by [mintlify](#)