



☰ Agent coordination > **Multi-Agent Sandbox & Tools**

Agent coordination

Multi-Agent Sandbox & Tools

Overview

Each agent in a multi-agent setup can now have its own:

Sandbox configuration (`agents.list[].sandbox` overrides `agents.defaults.sandbox`)

Tool restrictions (`tools.allow` / `tools.deny` , plus `agents.list[].tools`)

This allows you to run multiple agents with different security profiles:

Personal assistant with full access

Family/work agents with restricted tools

Public-facing agents in sandboxes

`setupCommand` belongs under `sandbox.docker` (global or per-agent) and runs once when the container is created.

Auth is per-agent: each agent reads from its own `agentDir` auth store at:

`~/.openclaw/agents/<agentId>/agent/auth-profiles.json`

Credentials are **not** shared between agents. Never reuse `agentDir` across agents. If you want to share creds, copy `auth-profiles.json` into the other agent's `agentDir`.

>

For how sandboxing behaves at runtime, see [**Sandboxing**](#). For debugging “why is this blocked?”, see [**Sandbox vs Tool Policy vs Elevated**](#) and `openclaw sandbox explain`.

Configuration Examples

Example 1: Personal + Restricted Family Agent



```
"agents": {
    "list": [
        {
            "id": "main",
            "default": true,
            "name": "Personal Assistant",
            "workspace": "~/.openclaw/workspace",
            "sandbox": { "mode": "off" }
        },
        {
            "id": "family",
            "name": "Family Bot",
            "workspace": "~/.openclaw/workspace-family",
            "sandbox": {
                "mode": "all",
                "scope": "agent"
            },
            "tools": {
                "allow": ["read"],
                "deny": ["exec", "write", "edit", "apply_patch", "process", "browser"]
            }
        }
    ]
},
"bindings": [
    {
        "agentId": "family",
        "match": {
            "provider": "whatsapp",
            "accountId": "*",
            "peer": {
                "kind": "group",
                "id": "120363424282127706@g.us"
            }
        }
    }
]
```

Result:



main agent: Runs on host, full tool access
family agent: Runs in Docker (one container per agent), only read tool

Example 2: Work Agent with Shared Sandbox

```
{  
  "agents": {  
    "list": [  
      {  
        "id": "personal",  
        "workspace": "~/.openclaw/workspace-personal",  
        "sandbox": { "mode": "off" }  
      },  
      {  
        "id": "work",  
        "workspace": "~/.openclaw/workspace-work",  
        "sandbox": {  
          "mode": "all",  
          "scope": "shared",  
          "workspaceRoot": "/tmp/work-sandboxes"  
        },  
        "tools": {  
          "allow": [ "read", "write", "apply_patch", "exec"],  
          "deny": [ "browser", "gateway", "discord"]  
        }  
      }  
    ]  
  }  
}
```

Example 2b: Global coding profile + messaging-only agent



```
{  
  "tools": { "profile": "coding" },  
  "agents": {  
    "list": [  
      {  
        "id": "support",  
        "tools": { "profile": "messaging", "allow": ["slack"] }  
      }  
    ]  
  }  
}
```

Result:

default agents get coding tools
support agent is messaging-only (+ Slack tool)

Example 3: Different Sandbox Modes per Agent



```
"agents": {  
    "defaults": {  
        "sandbox": {  
            "mode": "non-main", // Global default  
            "scope": "session"  
        }  
    },  
    "list": [  
        {  
            "id": "main",  
            "workspace": "~/.openclaw/workspace",  
            "sandbox": {  
                "mode": "off" // Override: main never sandboxed  
            }  
        },  
        {  
            "id": "public",  
            "workspace": "~/.openclaw/workspace-public",  
            "sandbox": {  
                "mode": "all", // Override: public always sandboxed  
                "scope": "agent"  
            },  
            "tools": {  
                "allow": ["read"],  
                "deny": ["exec", "write", "edit", "apply_patch"]  
            }  
        }  
    ]  
}
```

Configuration Precedence

When both global (`agents.defaults.*`) and agent-specific (`agents.list[].*`) configs exist:

Sandbox Config



Agent-specific settings override global:

```
>  
agents.list[].sandbox.mode > agents.defaults.sandbox.mode  
agents.list[].sandbox.scope > agents.defaults.sandbox.scope  
agents.list[].sandbox.workspaceRoot > agents.defaults.sandbox.workspaceRoot  
agents.list[].sandbox.workspaceAccess > agents.defaults.sandbox.workspaceAccess  
agents.list[].sandbox.docker.* > agents.defaults.sandbox.docker.*  
agents.list[].sandbox.browser.* > agents.defaults.sandbox.browser.*  
agents.list[].sandbox.prune.* > agents.defaults.sandbox.prune.*
```

Notes:

`agents.list[].sandbox.{docker,browser,prune}.*` overrides
`agents.defaults.sandbox.{docker,browser,prune}.*` for that agent (ignored
when sandbox scope resolves to "shared").

Tool Restrictions

The filtering order is:

1. **Tool profile** (`tools.profile` or `agents.list[].tools.profile`)
2. **Provider tool profile** (`tools.byProvider[provider].profile` or
`agents.list[].tools.byProvider[provider].profile`)
3. **Global tool policy** (`tools.allow` / `tools.deny`)
4. **Provider tool policy** (`tools.byProvider[provider].allow/deny`)
5. **Agent-specific tool policy** (`agents.list[].tools.allow/deny`)
6. **Agent provider policy**
(`agents.list[].tools.byProvider[provider].allow/deny`)
7. **Sandbox tool policy** (`tools.sandbox.tools` or
`agents.list[].tools.sandbox.tools`)
8. **Subagent tool policy** (`tools.subagents.tools`, if applicable)

Each level can further restrict tools, but cannot grant back denied tools from earlier levels. If `agents.list[].tools.sandbox.tools` is set, it replaces `tools.sandbox.tools` for that agent. If `agents.list[].tools.profile` is set, it overrides `tools.profile` for that agent. Provider tool keys accept either `provider` (e.g. `google-antigravity`) or `provider/model` (e.g. `openai/gpt-5.2`).

Tool groups (shorthands)

Tool policies (global, agent, sandbox) support `group:*` entries that expand to multiple concrete tools:

```
group:runtime : exec , bash , process  
group:fs : read , write , edit , apply_patch  
group:sessions : sessions_list , sessions_history , sessions_send ,  
sessions_spawn , session_status  
group:memory : memory_search , memory_get  
group:ui : browser , canvas  
group:automation : cron , gateway  
group:messaging : message  
group:nodes : nodes  
group:openclaw : all built-in OpenClaw tools (excludes provider  
plugins)
```

Elevated Mode

`tools.elevated` is the global baseline (sender-based allowlist). `agents.list[].tools.elevated` can further restrict elevated for specific agents (both must allow).

Mitigation patterns:

Deny `exec` for untrusted agents (`agents.list[].tools.deny: ["exec"]`)



Avoid allowlisting senders that route to restricted agents

Disable elevated globally (`tools.elevated.enabled: false`) if you only want sandboxed execution

>

Disable elevated per agent (`agents.list[].tools.elevated.enabled: false`) for sensitive profiles

Migration from Single Agent

Before (single agent):

```
{  
  "agents": {  
    "defaults": {  
      "workspace": "~/.openclaw/workspace",  
      "sandbox": {  
        "mode": "non-main"  
      }  
    }  
  },  
  "tools": {  
    "sandbox": {  
      "tools": {  
        "allow": ["read", "write", "apply_patch", "exec"],  
        "deny": []  
      }  
    }  
  }  
}
```

After (multi-agent with different profiles):



```
"agents": {  
    "list": [  
        {  
            "id": "main",  
            "default": true,  
            "workspace": "~/.openclaw/workspace",  
            "sandbox": { "mode": "off" }  
        }  
    ]  
}
```

Legacy `agent.*` configs are migrated by `openclaw doctor` ; prefer `agents.defaults + agents.list` going forward.

Tool Restriction Examples

Read-only Agent

```
{  
    "tools": {  
        "allow": ["read"],  
        "deny": ["exec", "write", "edit", "apply_patch", "process"]  
    }  
}
```

Safe Execution Agent (no file modifications)



```

"tools": {
  "allow": ["read", "exec", "process"],
  "deny": ["write", "edit", "apply_patch", "browser", "gateway"]
}
}

```

Communication-only Agent

```
{
  "tools": {
    "sessions": { "visibility": "tree" },
    "allow": ["sessions_list", "sessions_send", "sessions_history", "session_status"],
    "deny": ["exec", "write", "edit", "apply_patch", "read", "browser"]
  }
}
```

Common Pitfall: “non-main”

`agents.defaults.sandbox.mode: "non-main"` is based on `session.mainKey` (default `"main"`), not the agent id. Group/channel sessions always get their own keys, so they are treated as non-main and will be sandboxed. If you want an agent to never sandbox, set `agents.list[].sandbox.mode: "off"`.

Testing

After configuring multi-agent sandbox and tools:

1. Check agent resolution:



```
openclaw agents list --bindings  
>
```

2. Verify sandbox containers:

```
docker ps --filter "name=openclaw-sbx-"
```

3. Test tool restrictions:

Send a message requiring restricted tools

Verify the agent cannot use denied tools

4. Monitor logs:

```
tail -f "${OPENCLAW_STATE_DIR:-$HOME/.openclaw}/logs/gateway.log"
```

Troubleshooting

Agent not sandboxed despite mode: "all"

Check if there's a global `agents.defaults.sandbox.mode` that overrides it

Agent-specific config takes precedence, so set
`agents.list[].sandbox.mode: "all"`

Tools still available despite deny list

Check tool filtering order: global → agent → sandbox → subagent

Each level can only further restrict, not grant back



Verify with logs: [tools] filtering tools for agent:\${agentId}

Container not isolated per agent

Set scope: "agent" in agent-specific sandbox config

Default is "session" which creates one container per session

See Also

[Multi-Agent Routing](#)

[Sandbox Configuration](#)

[Session Management](#)

◀ Sub-Agents

Slash Commands ▶

Powered by [mintlify](#)