



## ☰ Agent coordination > Sub-Agents

Agent coordination

# Sub-Agents

Sub-agents are background agent runs spawned from an existing agent run. They run in their own session ( `agent:<agentId>:subagent:<uuid>` ) and, when finished, **announce** their result back to the requester chat channel.

## Slash command

Use `/subagents` to inspect or control sub-agent runs for the **current session**:

```
/subagents list  
  
/subagents kill <id|#|all>  
  
/subagents log <id|#> [limit] [tools]  
  
/subagents info <id|#>  
  
/subagents send <id|#> <message>  
  
/subagents steer <id|#> <message>  
  
/subagents spawn <agentId> <task> [--model <model>] [--thinking <level>]  
  
/subagents info shows run metadata (status, timestamps, session id, transcript path, cleanup).
```

## Spawn behavior

/subagents spawn starts a background sub-agent as a user command, not an internal relay, and it sends one final completion update back to the requester chat when the run finishes.

&gt;

The spawn command is non-blocking; it returns a run id immediately.

On completion, the sub-agent announces a summary/result message back to the requester chat channel.

For manual spawns, delivery is resilient:

OpenClaw tries direct agent delivery first with a stable idempotency key.

If direct delivery fails, it falls back to queue routing.

If queue routing is still not available, the announce is retried with a short exponential backoff before final give-up.

The completion message is a system message and includes:

Result ( assistant reply text, or latest toolResult if the assistant reply is empty)

Status ( completed successfully / failed / timed out )

compact runtime/token stats

--model and --thinking override defaults for that specific run.

Use info / log to inspect details and output after completion.

Primary goals:

Parallelize “research / long task / slow tool” work without blocking the main run.

Keep sub-agents isolated by default (session separation + optional sandboxing).

Keep the tool surface hard to misuse: sub-agents do **not** get session tools by default.

Support configurable nesting depth for orchestrator patterns.

Cost note: each sub-agent has its **own** context and token usage. For heavy or repetitive tasks, set a cheaper model for sub-agents and keep your main agent on a higher-quality model. You can configure this via `agents.defaults.subagents.model` or per-agent overrides.

## Tool

Use `sessions_spawn` :

Starts a sub-agent run (`deliver: false`, `global lane: subagent`)

Then runs an announce step and posts the announce reply to the requester chat channel

Default model: inherits the caller unless you set

`agents.defaults.subagents.model` (or per-agent

`agents.list[].subagents.model`); an explicit `sessions_spawn.model` still wins.

Default thinking: inherits the caller unless you set

`agents.defaults.subagents.thinking` (or per-agent

`agents.list[].subagents.thinking`); an explicit `sessions_spawn.thinking` still wins.

Tool params:

`task` (required)

`label?` (optional)

`agentId?` (optional; spawn under another agent id if allowed)

`model?` (optional; overrides the sub-agent model; invalid values are skipped and the sub-agent runs on the default model with a warning in the tool result)

`thinking?` (optional; overrides thinking level for the sub-agent run)

`runTimeoutSeconds?` (default 0 ; when set, the sub-agent run is aborted after N seconds)



cleanup? ( delete|keep , default keep )

#### Allowlist:

›

```
agents.list[].subagents.allowAgents : list of agent ids that can be  
targeted via agentId ( ["*"] to allow any). Default: only the  
requester agent.
```

#### Discovery:

```
Use agents_list to see which agent ids are currently allowed for  
sessions_spawn .
```

#### Auto-archive:

```
Sub-agent sessions are automatically archived after  
agents.defaults.subagents.archiveAfterMinutes (default: 60).
```

```
Archive uses sessions.delete and renames the transcript to  
*.deleted.<timestamp> (same folder).
```

```
cleanup: "delete" archives immediately after announce (still keeps  
the transcript via rename).
```

```
Auto-archive is best-effort; pending timers are lost if the gateway  
restarts.
```

```
runTimeoutSeconds does not auto-archive; it only stops the run. The  
session remains until auto-archive.
```

```
Auto-archive applies equally to depth-1 and depth-2 sessions.
```

## Nested Sub-Agents

By default, sub-agents cannot spawn their own sub-agents (`maxSpawnDepth: 1`). You can enable one level of nesting by setting `maxSpawnDepth: 2`, which allows the **orchestrator pattern**: main → orchestrator sub-agent → worker sub-sub-agents.

## How to enable



```
{
  agents: {
    defaults: {
      subagents: {
        maxSpawnDepth: 2, // allow sub-agents to spawn children (default: 1)
        maxChildrenPerAgent: 5, // max active children per agent session (default
        maxConcurrent: 8, // global concurrency lane cap (default: 8)
      },
    },
  },
}
```

## Depth levels

Depth	Session key shape	Role	Can spawn?
0	agent:<id>:main	Main agent	Always
1	agent:<id>:subagent:<uuid>	Sub-agent (orchestrator when depth 2 allowed)	Only if maxSpawnDepth >= 2
2	agent:<id>:subagent:<uuid>:subagent:<uuid>	Sub-sub-agent (leaf worker)	Never

## Announce chain

Results flow back up the chain:

1. Depth-2 worker finishes → announces to its parent (depth-1 orchestrator)
2. Depth-1 orchestrator receives the announce, synthesizes results, finishes → announces to main

3. Main agent receives the announce and delivers to the user



Each level only sees announces from its direct children.

>

## Tool policy by depth

**Depth 1 (orchestrator, when `maxSpawnDepth >= 2`)**: Gets `sessions_spawn`, `subagents`, `sessions_list`, `sessions_history` so it can manage its children. Other session/system tools remain denied.

**Depth 1 (leaf, when `maxSpawnDepth == 1`)**: No session tools (current default behavior).

**Depth 2 (leaf worker)**: No session tools – `sessions_spawn` is always denied at depth 2. Cannot spawn further children.

## Per-agent spawn limit

Each agent session (at any depth) can have at most `maxChildrenPerAgent` (default: 5) active children at a time. This prevents runaway fan-out from a single orchestrator.

## Cascade stop

Stopping a depth-1 orchestrator automatically stops all its depth-2 children:

`/stop` in the main chat stops all depth-1 agents and cascades to their depth-2 children.

`/subagents kill <id>` stops a specific sub-agent and cascades to its children.

`/subagents kill all` stops all sub-agents for the requester and cascades.

## Authentication

Sub-agent auth is resolved by **agent id**, not by session type:



The sub-agent session key is `agent:<agentId>:subagent:<uuid>`.

The auth store is loaded from that agent's `agentDir`.

The main agent's auth profiles are merged in as a **fallback**; agent profiles override main profiles on conflicts.

Note: the merge is additive, so main profiles are always available as fallbacks. Fully isolated auth per agent is not supported yet.

## Announce

Sub-agents report back via an announce step:

The announce step runs inside the sub-agent session (not the requester session).

If the sub-agent replies exactly `ANNOUNCE_SKIP`, nothing is posted.

Otherwise the announce reply is posted to the requester chat channel via a follow-up `agent` call (`deliver=true`).

Announce replies preserve thread/topic routing when available (Slack threads, Telegram topics, Matrix threads).

Announce messages are normalized to a stable template:

`Status`: derived from the run outcome (`success`, `error`, `timeout`, or `unknown`).

`Result`: the summary content from the announce step (or `(not available)` if missing).

`Notes`: error details and other useful context.

`Status` is not inferred from model output; it comes from runtime outcome signals.

Announce payloads include a stats line at the end (even when wrapped):



Runtime (e.g., `runtime 5m12s`)

Token usage (input/output/total)

Estimated cost when model pricing is configured

(`models.providers.*.models[].cost`)

`sessionKey`, `sessionId`, and transcript path (so the main agent can fetch history via `sessions_history` or inspect the file on disk)

## Tool Policy (sub-agent tools)

By default, sub-agents get **all tools except session tools** and system tools:

`sessions_list`

`sessions_history`

`sessions_send`

`sessions_spawn`

When `maxSpawnDepth >= 2`, depth-1 orchestrator sub-agents additionally receive `sessions_spawn`, `subagents`, `sessions_list`, and `sessions_history` so they can manage their children.

Override via config:



```
agents: {
  defaults: {
    subagents: { >
      maxConcurrent: 1,
    },
  },
},
tools: {
  subagents: {
    tools: {
      // deny wins
      deny: ["gateway", "cron"],
      // if allow is set, it becomes allow-only (deny still wins)
      // allow: ["read", "exec", "process"]
    },
  },
},
}
```

## Concurrency

Sub-agents use a dedicated in-process queue lane:

Lane name: `subagent`

Concurrency: `agents.defaults.subagents.maxConcurrent` (default 8)

## Stopping

Sending `/stop` in the requester chat aborts the requester session and stops any active sub-agent runs spawned from it, cascading to nested children.

`/subagents kill <id>` stops a specific sub-agent and cascades to its children.

# Limitations



Sub-agent announce is **best-effort**. If the gateway restarts, pending “announce back” work is lost.

Sub-agents still share the same gateway process resources; treat `maxConcurrent` as a safety valve.

```
sessions_spawn is always non-blocking: it returns { status:  
"accepted", runId, childSessionKey } immediately.
```

Sub-agent context only injects `AGENTS.md` + `TOOLS.md` (no `SOUL.md`, `IDENTITY.md`, `USER.md`, `HEARTBEAT.md`, or `BOOTSTRAP.md`).

Maximum nesting depth is 5 (`maxSpawnDepth` range: 1–5). Depth 2 is recommended for most use cases.

`maxChildrenPerAgent` caps active children per session (default: 5, range: 1–20).

[◀ Agent Send](#)

[Multi-Agent Sandbox & Tools ▶](#)

Powered by [mintlify](#)