



Overview

Tools

OpenClaw exposes **first-class agent tools** for browser, canvas, nodes, and cron. These replace the old `openclaw-*` skills: the tools are typed, no shelling, and the agent should rely on them directly.

Disabling tools

You can globally allow/deny tools via `tools.allow` / `tools.deny` in `openclaw.json` (deny wins). This prevents disallowed tools from being sent to model providers.

```
{  
  tools: { deny: ["browser"] },  
}
```

Notes:

Matching is case-insensitive.

* wildcards are supported ("*" means all tools).

If `tools.allow` only references unknown or unloaded plugin tool names, OpenClaw logs a warning and ignores the allowlist so core tools stay available.

Tool profiles (base allowlist)

 tools.profile sets a **base tool allowlist** before tools.allow / tools.deny .
Per-agent override: agents.list[].tools.profile .

Profiles: >

```
minimal : session_status only

coding : group:fs , group:runtime , group:sessions , group:memory ,
image

messaging : group:messaging , sessions_list , sessions_history ,
sessions_send , session_status

full : no restriction (same as unset)
```

Example (messaging-only by default, allow Slack + Discord tools too):

```
{
  tools: {
    profile: "messaging",
    allow: ["slack", "discord"],
  },
}
```

Example (coding profile, but deny exec/process everywhere):

```
{
  tools: {
    profile: "coding",
    deny: ["group:runtime"],
  },
}
```

Example (global coding profile, messaging-only support agent):



```

  tools: { profile: "coding" },
  agents: {
    list: [
      {
        id: "support",
        tools: { profile: "messaging", allow: ["slack"] },
      },
    ],
  },
}

```

Provider-specific tool policy

Use `tools.byProvider` to further restrict tools for specific providers (or a single `provider/model`) without changing your global defaults.

Per-agent override: `agents.list[].tools.byProvider`.

This is applied **after** the base tool profile and **before** allow/deny lists, so it can only narrow the tool set. Provider keys accept either `provider` (e.g. `google-antigravity`) or `provider/model` (e.g. `openai/gpt-5.2`).

Example (keep global coding profile, but minimal tools for Google Antigravity):

```
{
  tools: {
    profile: "coding",
    byProvider: {
      "google-antigravity": { profile: "minimal" },
    },
  },
}
```

Example (provider/model-specific allowlist for a flaky endpoint):

```
{
  tools: {
    allow: ["group:fs", "group:runtime", "sessions_list"],
    byProvider: {
      "openai/gpt-5.2": { allow: ["group:fs", "sessions_list"] },
    },
  },
}
```

Example (agent-specific override for a single provider):

```
{
  agents: {
    list: [
      {
        id: "support",
        tools: {
          byProvider: {
            "google-antigravity": { allow: ["message", "sessions_list"] },
          },
        },
      },
    ],
  },
}
```

Tool groups (shorthands)

Tool policies (global, agent, sandbox) support `group:*` entries that expand to multiple tools. Use these in `tools.allow` / `tools.deny`.

Available groups:

```
group:runtime : exec , bash , process
```



```

group:fs :  read ,  write ,  edit ,  apply_patch
group:sessions :  sessions_list ,  sessions_history ,  sessions_send ,
sessions_spawn ,  session_status
>
group:memory :  memory_search ,  memory_get
group:web :  web_search ,  web_fetch
group:ui :  browser ,  canvas
group:automation :  cron ,  gateway
group:messaging :  message
group:nodes :  nodes
group:openclaw : all built-in OpenClaw tools (excludes provider
plugins)

```

Example (allow only file tools + browser):

```
{
  tools: {
    allow: ["group:fs", "browser"],
  },
}
```

Plugins + tools

Plugins can register **additional tools** (and CLI commands) beyond the core set. See [Install + config](#) for install + config, and [Tool usage](#) for how tool usage guidance is injected into prompts. Some plugins ship their own skills alongside tools (for example, the voice-call plugin).

Optional plugin tools:

- `: typed workflow runtime with resumable approvals (requires the Lobster CLI on the gateway host).`



LLM Task: JSON-only LLM step for structured workflow output (optional schema validation).

>

Tool inventory

apply_patch

Apply structured patches across one or more files. Use for multi-hunk edits. Experimental: enable via `tools.exec.applyPatch.enabled` (OpenAI models only). `tools.exec.applyPatch.workspaceOnly` defaults to `true` (workspace-contained). Set it to `false` only if you intentionally want `apply_patch` to write/delete outside the workspace directory.

exec

Run shell commands in the workspace.

Core parameters:

`command` (required)

`yieldMs` (auto-background after timeout, default 10000)

`background` (immediate background)

`timeout` (seconds; kills the process if exceeded, default 1800)

`elevated` (bool; run on host if elevated mode is enabled/allowed; only changes behavior when the agent is sandboxed)

`host` (`sandbox` | `gateway` | `node`)

`security` (`deny` | `allowlist` | `full`)

`ask` (`off` | `on-miss` | `always`)

`node` (node id/name for `host=node`)

Need a real TTY? Set `pty: true`.

Notes:



Returns status: "running" with a `sessionId` when backgrounded.

Use `process` to poll/log/write/kill/clear background sessions.

If `process` is disallowed, `exec` runs synchronously and ignores `yieldMs` / `background`.

`elevated` is gated by `tools.elevated` plus any `agents.list[].tools.elevated` override (both must allow) and is an alias for `host=gateway + security=full`.

`elevated` only changes behavior when the agent is sandboxed (otherwise it's a no-op).

`host=node` can target a macOS companion app or a headless node host (`openclaw node run`).

gateway/node approvals and allowlists: [Exec approvals](#).

process

Manage background exec sessions.

Core actions:

`list` , `poll` , `log` , `write` , `kill` , `clear` , `remove`

Notes:

`poll` returns new output and exit status when complete.

`log` supports line-based `offset` / `limit` (omit `offset` to grab the last N lines).

`process` is scoped per agent; sessions from other agents are not visible.

loop-detection (tool-call loop guardrails)

OpenClaw tracks recent tool-call history and blocks or warns when it detects repetitive no-progress loops. Enable with

 tools.loopDetection.enabled: true (default is false).

```
{
  tools: {
    loopDetection: {
      enabled: true,
      warningThreshold: 10,
      criticalThreshold: 20,
      globalCircuitBreakerThreshold: 30,
      historySize: 30,
      detectors: {
        genericRepeat: true,
        knownPollNoProgress: true,
        pingPong: true,
      },
    },
  },
}
```

genericRepeat : repeated same tool + same params call pattern.

knownPollNoProgress : repeating poll-like tools with identical outputs.

pingPong : alternating A/B/A/B no-progress patterns.

Per-agent override: agents.list[].tools.loopDetection .

web_search

Search the web using Brave Search API.

Core parameters:

query (required)

count (1-10; default from tools.web.search.maxResults)

Notes:



Requires a Brave API key (recommended: `openclaw configure --section web`, or set `BRAVE_API_KEY`).

Enable via `tools.web.search.enabled`.

>

Responses are cached (default 15 min).

See [Web tools](#) for setup.

web_fetch

Fetch and extract readable content from a URL (HTML → markdown/text).

Core parameters:

`url` (required)

`extractMode` (`markdown` | `text`)

`maxChars` (truncate long pages)

Notes:

Enable via `tools.web.fetch.enabled`.

`maxChars` is clamped by `tools.web.fetch.maxCharsCap` (default 50000).

Responses are cached (default 15 min).

For JS-heavy sites, prefer the browser tool.

See [Web tools](#) for setup.

See [Firecrawl](#) for the optional anti-bot fallback.

browser

Control the dedicated OpenClaw-managed browser.

Core actions:

`status` , `start` , `stop` , `tabs` , `open` , `focus` , `close`

`snapshot` (aria/ai)



screenshot (returns image block + MEDIA:<path>)

act (UI actions:

click/type/press/hover/drag/select/fill/resize/wait/evaluate)

>

navigate , console , pdf , upload , dialog

Profile management:

profiles – list all browser profiles with status

create-profile – create new profile with auto-allocated port (or cdpUrl)

delete-profile – stop browser, delete user data, remove from config (local only)

reset-profile – kill orphan process on profile's port (local only)

Common parameters:

profile (optional; defaults to browser.defaultProfile)

target (sandbox | host | node)

node (optional; picks a specific node id/name) Notes:

Requires browser.enabled=true (default is true ; set false to disable).

All actions accept optional profile parameter for multi-instance support.

When profile is omitted, uses browser.defaultProfile (defaults to "chrome").

Profile names: lowercase alphanumeric + hyphens only (max 64 chars).

Port range: 18800-18899 (~100 profiles max).

Remote profiles are attach-only (no start/stop/reset).

If a browser-capable node is connected, the tool may auto-route to it (unless you pin target).



`snapshot` defaults to `ai` when Playwright is installed; use `aria` for the accessibility tree.

`snapshot` also supports role-snapshot options (`interactive` , `compact` , `depth` , `selector`) which return refs like `e12` .

`act` requires `ref` from `snapshot` (numeric `12` from AI snapshots, or `e12` from role snapshots); use `evaluate` for rare CSS selector needs.

Avoid `act` → `wait` by default; use it only in exceptional cases (no reliable UI state to wait on).

`upload` can optionally pass a `ref` to auto-click after arming.

`upload` also supports `inputRef` (aria ref) or `element` (CSS selector) to set `<input type="file">` directly.

canvas

Drive the node Canvas (present, eval, snapshot, A2UI).

Core actions:

```
present , hide , navigate , eval  
snapshot (returns image block + MEDIA:<path> )  
a2ui_push , a2ui_reset
```

Notes:

Uses gateway `node.invoke` under the hood.

If no `node` is provided, the tool picks a default (single connected node or local mac node).

A2UI is v0.8 only (no `createSurface`); the CLI rejects v0.9 JSONL with line errors.

Quick smoke: `openclaw nodes canvas a2ui push --node <id> --text "Hello from A2UI"` .

nodes



Discover and target paired nodes; send notifications; capture camera/screen.

Core actions:

```
status , describe
pending , approve , reject (pairing)
notify (macOS system.notify )
run (macOS system.run )
camera_snap , camera_clip , screen_record
location_get
```

Notes:

Camera/screen commands require the node app to be foregrounded.

Images return image blocks + MEDIA:<path> .

Videos return FILE:<path> (mp4) .

Location returns a JSON payload (lat/lon/accuracy/timestamp) .

run params: command argv array; optional cwd , env (KEY=VAL) , commandTimeoutMs , invokeTimeoutMs , needsScreenRecording .

Example (run):

```
{
  "action": "run",
  "node": "office-mac",
  "command": ["echo", "Hello"],
  "env": ["FOO=bar"],
  "commandTimeoutMs": 12000,
  "invokeTimeoutMs": 45000,
  "needsScreenRecording": false
}
```

image



Analyze an image with the configured image model.

Core parameters:

`image` (required path or URL)

`prompt` (optional; defaults to “Describe the image.”)

`model` (optional override)

`maxBytesMb` (optional size cap)

Notes:

Only available when `agents.defaults.imageModel` is configured (primary or fallbacks), or when an implicit image model can be inferred from your default model + configured auth (best-effort pairing).

Uses the image model directly (independent of the main chat model).

message

Send messages and channel actions across Discord/Google Chat/Slack/Telegram/WhatsApp/Signal/iMessage/MS Teams.

Core actions:

`send` (text + optional media; MS Teams also supports `card` for Adaptive Cards)

`poll` (WhatsApp/Discord/MS Teams polls)

`react` / `reactions` / `read` / `edit` / `delete`

`pin` / `unpin` / `list-pins`

`permissions`

`thread-create` / `thread-list` / `thread-reply`

`search`

`sticker`



```

member-info / role-info
emoji-list / emoji-upload / sticker-upload
role-add, / role-remove
channel-info / channel-list
voice-status
event-list / event-create
timeout / kick / ban

```

Notes:

`send` routes WhatsApp via the Gateway; other channels go direct.

`poll` uses the Gateway for WhatsApp and MS Teams; Discord polls go direct.

When a message tool call is bound to an active chat session, sends are constrained to that session's target to avoid cross-context leaks.

cron

Manage Gateway cron jobs and wakeups.

Core actions:

```

status , list
add , update , remove , run , runs
wake (enqueue system event + optional immediate heartbeat)

```

Notes:

`add` expects a full cron job object (same schema as `cron.add` RPC).
`update` uses `{ jobId, patch }` (`id` accepted for compatibility).

gateway

Restart or apply updates to the running Gateway process (in-place).

Core actions:

```
restart (authorizes + sends SIGUSR1 for in-process restart;
openclaw gateway restart in-place)

config.get / config.schema

config.apply (validate + write config + restart + wake)

config.patch (merge partial update + restart + wake)

update.run (run update + restart + wake)
```

Notes:

Use `delayMs` (defaults to 2000) to avoid interrupting an in-flight reply.

`restart` is disabled by default; enable with `commands.restart: true`.

sessions_list / sessions_history / sessions_send / sessions_spawn / session_status

List sessions, inspect transcript history, or send to another session.

Core parameters:

```
sessions_list : kinds? , limit? , activeMinutes? , messageLimit? (0 = none)

sessions_history : sessionKey (or sessionId) , limit? , includeTools?

sessions_send : sessionKey (or sessionId) , message , timeoutSeconds? (0 = fire-and-forget)

sessions_spawn : task , label? , agentId? , model? ,
runTimeoutSeconds? , cleanup?
```



```
session_status : sessionKey? (default current; accepts sessionId ),  
model? ( default clears override)
```

Notes: >

`main` is the canonical direct-chat key; global/unknown are hidden.

`messageLimit > 0` fetches last N messages per session (tool messages filtered).

Session targeting is controlled by `tools.sessions.visibility` (default tree : current session + spawned subagent sessions). If you run a shared agent for multiple users, consider setting `tools.sessions.visibility: "self"` to prevent cross-session browsing.

`sessions_send` waits for final completion when `timeoutSeconds > 0`.

Delivery/announce happens after completion and is best-effort; `status: "ok"` confirms the agent run finished, not that the announce was delivered.

`sessions_spawn` starts a sub-agent run and posts an announce reply back to the requester chat.

Reply format includes `Status`, `Result`, and compact stats.

`Result` is the assistant completion text; if missing, the latest `toolResult` is used as fallback.

Manual completion-mode spawns send directly first, with queue fallback and retry on transient failures (`status: "ok"` means run finished, not that announce delivered).

`sessions_spawn` is non-blocking and returns `status: "accepted"` immediately.

`sessions_send` runs a reply-back ping-pong (reply `REPLY_SKIP` to stop; max turns via `session.agentToAgent.maxPingPongTurns`, 0-5).

After the ping-pong, the target agent runs an `announce step`; reply `ANNOUNCE_SKIP` to suppress the announcement.

 Sandbox clamp: when the current session is sandboxed and
agents.defaults.sandbox.sessionToolsVisibility: "spawned" , OpenClaw clamps
tools.sessions.visibility to tree .

>

agents_list

List agent ids that the current session may target with sessions_spawn .

Notes:

Result is restricted to per-agent allowlists
(agents.list[].subagents.allowAgents).

When ["*"] is configured, the tool includes all configured agents
and marks allowAny: true .

Parameters (common)

Gateway-backed tools (canvas , nodes , cron):

```
gatewayUrl (default ws://127.0.0.1:18789 )
gatewayToken (if auth enabled)
timeoutMs
```

Note: when gatewayUrl is set, include gatewayToken explicitly. Tools do not inherit config or environment credentials for overrides, and missing explicit credentials is an error.

Browser tool:

```
profile (optional; defaults to browser.defaultProfile )
target ( sandbox | host | node )
node (optional; pin a specific node id/name)
```

Recommended agent flows



Browser automation:

1. browser → status / start
2. snapshot (ai or aria)
3. act (click/type/press)
4. screenshot if you need visual confirmation

Canvas render:

1. canvas → present
2. a2ui_push (optional)
3. snapshot

Node targeting:

1. nodes → status
2. describe on the chosen node
3. notify / run / camera_snap / screen_record

Safety

Avoid direct system.run ; use nodes → run only with explicit user consent.

Respect user consent for camera/screen capture.

Use status/describe to ensure permissions before invoking media commands.

How tools are presented to the agent

Tools are exposed in two parallel channels:

1. **System prompt text:** a human-readable list + guidance.

 2. **Tool schema:** the structured function definitions sent to the model API.
-

That means the agent sees both “what tools exist” and “how to call them.” If a tool doesn’t appear in the system prompt or the schema, the model cannot call it.

[Lobster >](#)

Powered by [mintlify](#)